Exogenous Features in Intraday Forecasts of Closing-Half-Hour Variance

Bullet-point draft of the full paper

James Chen Mihai Cucuringu Christopher S. Jones Austin Pollok

Working draft, 18 October 2026 — not for citation

Abstract

We investigate whether exogenous predictors improve intraday forecasts of a futures-based (E-mini) proxy for the realized variance of the S&P 500 over the closing half-hour interval (15:30–16:00 ET). Specifically, we estimate pooled models on the 13 regular-hours half-hour bars of each trading day with a rolling window of 1,000 trading days and evaluate the forecast issued at 15:30 for the closing half hour over the out-of-sample period from 3 January 2020 to 30 April 2024 ($T_{\text{oos}} = 1{,}082$ trading days), and the forecasts of every regular-hours bar from 3 January 2012 to 30 April 2024 (3,075 trading days). We compare a benchmark heterogeneous autoregressive (HAR) predictor set against an augmented predictor set that incorporates exogenous series. Under the QLIKE loss, the full predictor set yields a lower average out-of-sample loss than the HAR benchmark in 4 of the 6 intraday model classes, with 0 of these reductions statistically significant at the 5% level. Furthermore, although every intraday specification remains less accurate under QLIKE than a HAR-type model in the spirit of Corsi (2009), estimated by OLS on the closing-interval observations of the target's trailing averages over $k \in \{1, 5, 25, 125, 625, 3125\}$ samples plus calendar columns, incorporating the exogenous predictors narrows this gap in 4 of the 6 model classes. In an economic application to zero-days-to-expiration (0DTE) S&P 500 straddles on expiration days, a sign-based trading strategy achieves a higher Sharpe ratio under the full predictor set than under the HAR benchmark across every intraday model class.

Notes. Colored notes are draft comments, not paper text.
Red: two live statements disagree.
Orange: a sentence does not match the computation behind a printed number.
Blue: a reader cannot recover the input or the rule.
The first line is the decision. The second line is the problem. The rest is the exposition.

E34. Decision: base 5 across the paper (user, 2026-10-19); the pooled trees are the base-5 refit (user, 2026-10-22); the six pooled linear models of Sections 6 and 7 are the forecasts of Table 1 on the same design (user, 2026-10-22), so the 12 pooled intraday forecasts share one design: the 13 regular-hours bars, a rolling window of 1,000 trading days, evaluation from 3 January 2012.
Problem: (i) the model set (the 112 forecasts of the master table) still holds earlier designs: the earlier draft's pooled 48-interval forecasts (the "paper" family, among them its LightGBM and XGBoost) and the master table's pooled linear runs on all 48 intervals with $W = 96{,}000$ intervals (marked "All 48 intervals" in generated/rehaul_results_t3_full.tex, not in the paper), so the statistics computed over the model set in Section 7 and in the appendix (the counts against the short straddle, the sizing and compounding medians, the interval-shift test) include them. (ii) On the one design, the full predictor set lowers the closing-interval QLIKE in 4 of the 6 pooled classes (0 significantly) and raises the sign strategy's Sharpe ratio in 6 of them (note N15c); the sentences that said otherwise now print counts (writeup/E34_OVERNIGHT_REPORT.md). Table 4 gives the shared window.
The pooled trees (LightGBM, XGBoost and the random forest, on the HAR benchmark and on the full predictor set) are the refit on the design of the pooled linear models of Table 1:
results/pooled_trees_rerun13_2026-10-20/w1000 (CARC; branch insurance-selling-2026-09-19-trees-results, commit 6746103), base-5 HAR horizons, FOMC series in, the frozen configurations, re-estimated at every bar on a rolling window of 13,000 bars, from 3 January 2012 to 30 April 2024. writeup/pooled_trees_13.py reads their 16:00 rows for make_rehaul_results.py, make_rehaul_economics.py and make_rehaul_oos_table.py, and writeup/make_rehaul_oos_rerun.py scores every bar of them in Table 1 (note E36), so the counts in this abstract and in the introduction (6 classes, 0 significant), the pooled comparisons of Sections 6 and 7 and of the appendix now use them; the tree rows of Table 3 are closing-interval models and are unchanged. The intraday OLS of the earlier Table 1 left the table with note E36. With the refit, the QLIKE reduction from the full predictor set is statistically significant at the 5% level in 0 of the three tree classes (all three before), and the sentences that said each tree difference was significant now print counts (writeup/E34_OVERNIGHT_REPORT.md).

1 Introduction

- **Research question.** We investigate whether incorporating exogenous predictors improves intraday forecasts of a futures-based (E-mini) proxy for the realized variance of the S&P 500 over the closing half-hour interval.
  - Specifically, at the 15:30 ET forecast origin on trading day t, we seek to predict RV_t, which denotes the realized variance over the 15:30–16:00 closing interval (see Section 3).
  - Each pooled intraday model is estimated on the 13 regular-hours half-hour bars of the trading day, and we evaluate its out-of-sample predictions on the forecast issued at 15:30 for the closing half hour.
  - To isolate the contribution of the additional covariates, we compare the full predictor set against the HAR benchmark; the exogenous series constitute the exact difference between the two sets (see Section 4).
- **Motivation for the HAR baseline.** Following Corsi (2009), the standard HAR regression models realized volatility solely as a linear function of its own past lags. Consequently, it omits a rich set of exogenous variables available by 15:30, including higher-order realized moments, liquidity measures, the cross-section of constituent returns, news and social-media sentiment, option-implied volatility, volatility demand indicators, and scheduled macroeconomic and earnings announcements.
- **Empirical design.** For each intraday interval (t, i), we consider a single forecast mapping of the form
$$\hat{y}_{t,i} = f_{m,\mathcal{X},\mathcal{I}}(\boldsymbol{x}_{t,i}),$$
indexed by the model class m and the predictor set $\mathcal{X} \in \{\mathrm{HAR}, \mathrm{Full}\}$.
  - The primary empirical results reported in Section 6 focus on the pooled intraday estimation scheme, $\mathcal{I} = \{1, \ldots, M\}$.

  - As our primary benchmark forecast, we employ the ordinary least squares (OLS) regression on the HAR benchmark predictor set, estimated exclusively on the closing interval (the closing-interval HAR benchmark; see Sections 5.2 and 5.3).
  - We evaluate predictive accuracy via the average QLIKE loss over the out-of-sample evaluation span $\mathcal{T}_{\text{oos}}$ from 3 January 2020 to 30 April 2024 ($T_{\text{oos}} = 1{,}082$ trading days; see Section 5.3); the accuracy of the pooled linear models and tree ensembles in Table 1 is scored on every regular-hours bar from 3 January 2012 to 30 April 2024 (3,075 trading days, a rolling window of 1,000 trading days).
- **Main findings.**
  - Incorporating the full predictor set reduces the average out-of-sample QLIKE loss relative to the HAR benchmark in 4 of the 6 intraday model classes, with 0 of these reductions statistically significant at the 5% level (see Section 6). The reduction is statistically significant at the 5% level in 0 of the three linear classes and in 0 of the three tree classes.
  - Every pooled intraday forecast incurs a higher average QLIKE loss than the closing-interval HAR benchmark. Under the full predictor set, this performance gap is statistically significant at the 5% level for 3 of the three linear models and for 1 of the three tree ensembles. Augmenting the intraday models with the full predictor set narrows this gap in 4 of the 6 model classes (see Section 6.4).
- **Economic application.** To assess the economic value of the forecasts on trading days with a same-day SPXW expiration, we construct a trading strategy that takes a long or short position in a zero-days-to-expiration (0DTE) at-the-money straddle based on the sign of the difference between the variance forecast and option-implied variance (see Section 7). Using the full predictor set increases the annualized Sharpe ratio of this strategy relative to the HAR benchmark in every intraday model class (see Section 7.3).

  N15c. Problem: under the Sharpe ratio of note D6, the full predictor set raises the pooled Sharpe ratio in 6 of the 6 classes (ΔSR from +0.23 to +0.97), so "every intraday model class" holds in point estimates; 0 of these gains are statistically significant at the 5% level (bootstrap test), and under the previous ratio the count was 2 of 6. See notes N15c and D6 in Section 7.
- **Paper organization.** The remainder of the paper is organized as follows. Section 2 reviews the related literature. Section 3 describes the data and target construction, and Section 4 details the benchmark and exogenous predictors. Section 5 presents the forecasting models and rolling estimation protocol. Section 6 reports the out-of-sample forecasting results. Section 7 evaluates the economic value of the forecasts in the 0DTE straddle strategy, and Section 8 concludes.

2 Related work

- **Realized variance and the HAR model.** Following Corsi (2009), the heterogeneous autoregressive (HAR) model specifies the conditional expectation of next-day realized variance as a linear function of past realized variances aggregated over daily, weekly, and monthly horizons: $\mathbb{E}_t[\mathrm{RV}_{t+1}^{(d)}] = \beta_0 + \beta_d \,\mathrm{RV}_t^{(d)} + \beta_w \,\mathrm{RV}_t^{(w)} + \beta_m \,\mathrm{RV}_t^{(m)}$, where $\mathbb{E}_t$ denotes the conditional expectation given information available at the forecast origin on day t, and $\mathrm{RV}_t^{(d)}$, $\mathrm{RV}_t^{(w)}$, and $\mathrm{RV}_t^{(m)}$ denote the average daily realized variance over the preceding 1, 5, and 22 trading days, respectively. Because all three components represent daily realized variances aggregated at the daily, weekly, and monthly horizons, forecasting the realized variance of a specific intraday

interval lies outside this standard daily specification.
See Corsi (2009), Andersen et al. (2003), Barndorff-Nielsen and Shephard (2004), Andersen et al. (2007), Corsi et al. (2008), Bollerslev et al. (2016), Bekierman and Manner (2018), Bennedsen et al. (2022), Hansen et al. (2012), and Andersen et al. (2006).

  - In a multivariate setting, the panel HAR framework of Bollerslev et al. (2018) estimates a single pooled regression jointly across a cross-section of assets.
  - For the theoretical foundations of realized variance, see Barndorff-Nielsen and Shephard (2002).

- **Diurnal pattern.** It is well documented that intraday return volatility exhibits a pronounced U-shaped pattern across the trading session. To account for this periodicity, standard practice scales intraday realized variance by a diurnal factor estimated from the corresponding clock-time interval on prior trading days.
  See Andersen and Bollerslev (1997) on intraday periodicity, Engle and Gallo (2006) on intraday range and volatility measures as indicators of daily volatility, and Hizmeri and Bevilacqua (2025) on morning volatility uncertainty.
  - See also Taylor and Xu (1997) on diurnal volatility adjustments and Engle and Sokalska (2012) on multiplicative decompositions of intraday volatility.
- **0DTE options and the variance risk premium.** The variance risk premium, defined as the difference between option-implied variance and expected realized variance, $\mathbb{E}_t^{\mathbb{Q}}[\mathrm{RV}_t] - \mathbb{E}_t^{\mathbb{P}}[\mathrm{RV}_t]$, under the risk-neutral ($\mathbb{Q}$) and physical ($\mathbb{P}$) measures (where RV_t denotes the realized variance through expiration), is empirically positive on average.
  See Hizmeri and Bevilacqua (2025), as well as the foundational option-pricing frameworks of Black and Scholes (1973) and Merton (1973).
  - On zero-days-to-expiration (0DTE) options, see Bandi et al. (2023), Beckmeyer et al. (2023), Brogaard et al. (2023), Dim et al. (2024), Vilkov (2026), and Almeida et al. (2025); on the variance risk premium more broadly, see Carr and Wu (2009) and Bollerslev et al. (2009).
- **Machine learning for volatility and returns.** A growing literature models conditional return and volatility dynamics via flexible nonlinear mappings $f(\boldsymbol{x}_t)$ of a high-dimensional predictor vector $\boldsymbol{x}_t$.
  See Gu et al. (2020) for return prediction and Christensen et al. (2023) for realized volatility forecasting; for gradient-boosted decision trees, see Ke et al. (2017) (LightGBM) and Chen and Guestrin (2016) (XGBoost); for Bayesian hyperparameter optimization, see Akiba et al. (2019) (Optuna).
  - For ensemble methods, see Breiman (2001) on random forests and Friedman (2001) on gradient boosting; for machine-learning applications to realized volatility and intraday commonality, see Bucci (2020), Rahimikia and Poon (2020), and Zhang et al. (2024).
- **Shrinkage with many weak signals.** In high-dimensional linear regressions, penalized estimators regularize the coefficient vector via

$$\hat{\boldsymbol{\beta}} = \arg\min_{\boldsymbol{\beta}} \|\boldsymbol{y} - \boldsymbol{X}\boldsymbol{\beta}\|_2^2 + \lambda(\alpha\|\boldsymbol{\beta}\|_1 + \tfrac{1-\alpha}{2}\|\boldsymbol{\beta}\|_2^2)$$

  for a response vector $\boldsymbol{y}$, a predictor matrix $\boldsymbol{X}$, and an overall penalty strength λ, nesting ridge regression ($\alpha = 0$) and the lasso ($\alpha = 1$) as special cases; multi-penalty ridge extensions assign distinct regularization weights across feature blocks via $\lambda \sum_j \omega_j \beta_j^2$, where the penalty weight ω_j is held constant within each block.

See Giannone et al. (2021) on dense versus sparse representations, Shen and Xiu (2025) on learning from many weak signals, Kozak et al. (2020) on shrinking the cross-section of returns, Hoerl and Kennard (1970) on ridge regression, Tibshirani (1996) on the lasso, Zou and Hastie (2005) on the elastic net, and van de Wiel et al. (2021) and van Wieringen (2023) on multi-penalty ridge estimation.

  - See also Zou (2006) on the adaptive lasso and feature-specific penalty weights, and Friedman et al. (2010) on coordinate-descent estimation under the elastic-net parametrization.

3 Data

- **Data sources.** We group the exogenous covariates into eight distinct categories. Specifically, the pooled intraday specifications employ the 46 series across all categories, the FOMC group included. A complete list and definition of each individual series is provided in Appendix Table 6.
  - **Realized moments** (7 series). We obtain intraday data on the front-month E-mini S&P 500 futures contract (defined as the most actively traded contract), which trades nearly 24 hours a day from 1998 through 2024, aggregated onto a 30-minute grid. Here, $t = 1, \ldots, T$ indexes trading days and $i = 1, \ldots, M$ indexes the half-hour intervals within day t, each ending on a half-hour timestamp; in particular, $i = c$ denotes the 15:30–16:00 ET closing interval. Each realized moment for interval (t, i) is computed by summing the corresponding one-minute quantity across all non-missing minutes within that half hour.
  - Let $r_{t,i,k}$ denote the one-minute log-price change of the front-month contract for the k-th non-missing minute in interval i of day t, where $k = 1, \ldots, N_{t,i}$ and $N_{t,i} \leq 30$. Following Andersen et al. (2003), we define the realized variance over interval (t, i) as

$$\mathrm{RV}_{t,i} = \sum_{k=1}^{N_{t,i}} r_{t,i,k}^2.$$

  - We construct the 7 exogenous realized-moment series as analogous one-minute sums over interval (t, i) (note that $\mathrm{RV}_{t,i}$ above serves as our forecast target and is not counted among the exogenous series):

$$\texttt{sumret}_{t,i} = \sum_k r_{t,i,k}, \quad \texttt{sumabsret}_{t,i} = \sum_k |r_{t,i,k}|, \quad \texttt{sumret3}_{t,i} = \sum_k r_{t,i,k}^3,$$
$$\texttt{sumret4}_{t,i} = \sum_k r_{t,i,k}^4, \quad \texttt{sumpret2}_{t,i} = \sum_k r_{t,i,k}^2 \, \mathbf{1}\{r_{t,i,k} > 0\},$$
$$\texttt{sumbipow}_{t,i} = \sum_{k \geq 2} |r_{t,i,k}| \, |r_{t,i,k-1}|, \quad \texttt{sumautocov}_{t,i} = \sum_{k \geq 2} r_{t,i,k} \, r_{t,i,k-1}.$$

    Specifically, the third sum corresponds to the unnormalized realized skewness of Amaya et al. (2015) (that is, prior to dividing by $\mathrm{RV}_{t,i}^{3/2}/\sqrt{N_{t,i}}$), and the fourth sum is the unscaled realized quarticity (without the factor $N_{t,i}/3$). Furthermore, $\texttt{sumpret2}_{t,i}$ denotes the positive realized semivariance of Barndorff-Nielsen et al. (2010), $\texttt{sumbipow}_{t,i}$ denotes the unscaled bipower variation of Barndorff-Nielsen and Shephard (2004) (without the $\pi/2$ normalization factor), and $\texttt{sumautocov}_{t,i}$ denotes the first-order realized autocovariance.

– **Liquidity** (14 series). From the front-month futures contract, we include total contract volume traded during the half hour as well as the number of observed one-minute intervals $N_{t,i}$ (`numobs`). In addition, using intraday data on S&P 500 constituents from 1993 to 2024, we construct both equal-weighted and value-weighted cross-sectional aggregates of share turnover, buyer-initiated turnover, seller-initiated turnover, the effective spread, the quoted spread, and order-flow imbalance. Here, constituent turnover is defined as the number of shares traded during the half hour divided by prior-day shares outstanding, while buyer- and seller-initiated turnover aggregate trades according to their estimated trade direction. The effective spread measures the average percentage deviation between the transaction price and the prevailing mid-quote, and the quoted spread measures the average percentage bid–ask spread. Equal-weighted aggregates take simple cross-sectional averages across constituents, whereas value-weighted aggregates weight each constituent by its prior-day market capitalization. For each weighting scheme (`ofi_ewstock` and `ofi_vwstock`), we define order-flow imbalance as

$$\mathrm{OFI}_{t,i} = \frac{\mathrm{Buy}_{t,i} - \mathrm{Sell}_{t,i}}{\mathrm{Buy}_{t,i} + \mathrm{Sell}_{t,i}} \in [-1, 1], \qquad \text{missing if } \mathrm{Buy}_{t,i} + \mathrm{Sell}_{t,i} = 0,$$

where $\mathrm{Buy}_{t,i}$ and $\mathrm{Sell}_{t,i}$ denote the aggregate buyer- and seller-initiated turnover in interval i of day t, respectively.

– **Equal-weighted market** (6 series) and **value-weighted market** (6 series). For each weighting scheme described above, we compute six cross-sectional constituent return moments: realized variance, realized skewness, realized quarticity, absolute variation, bipower variation, and positive semivariance. In contrast to the futures moments, which are constructed from one-minute log-price changes, these constituent moments are computed from 5-minute log returns. Moreover, the constituent moment groups omit signed returns and realized autocovariance, whereas the exogenous futures moment group omits realized variance (which serves as the target variable).

– **Sentiment** (3 series). Using StockTwits message streams beginning in 2009 for total posts and 2010 for user-tagged posts, we construct three half-hourly indicators: investor attention, defined as the total number of posts within the half hour; a sentiment score that equals 1 when all tagged posts are bullish and -1 when all tagged posts are bearish; and the sentiment count, defined as the number of posts carrying either a bullish or bearish tag.

E30. Decision: open, for the professor.
Problem: the text gives only the ± 1 endpoints of the sentiment score.
Question for the professor: what is the formula of the score between the endpoints, and what value does it take on intervals with no tagged posts? The repository receives the score precomputed (data/spy_and_sentiment.parquet, column stocktwits_sentiment, first observed in June 2010); src/data/loading.py:280 reads it as a net score in $[-1, 1]$, and no script in src/ or experiments/ computes it.

– **Implied volatility** (3 series). We include three Cboe volatility indices: the 30-day S&P 500 implied volatility index (VIX, available from 2006), the three-month implied volatility index (VIX3M, from 2009), and the implied volatility of the VIX (VVIX, from 2012), all extending through February 2024.

– **Volatility demand** (4 series). Using Cboe option-volume records from 2011 through August 2023, we measure net vega purchased by non-market-maker participants across four series formed by crossing two option universes—options on the S&P 500 index and top S&P 500 exchange-traded funds versus all exchange-traded options—with two

transaction subsets, namely all opening and closing customer trades versus opening customer trades only.

– **FOMC** (3 series; used by the pooled intraday model and by the closing-interval specifications of Appendix G). Based on scheduled Federal Open Market Committee announcement timestamps from 1993 through November 2023, we define 3 series: an indicator for whether an announcement occurs on day t and two proximity measures $1/(1+n^{+}_{t,i})$ and $1/(1+n^{-}_{t,i})$, where $n^{+}_{t,i}$ and $n^{-}_{t,i}$ denote the number of intervals until the next scheduled announcement and since the most recent scheduled announcement, respectively. Because the meeting schedule is announced in advance, these variables are known prior to the start of the interval.

E28. Decision (user, 2026-10-22): disclose the tails in the text; add the flag UNOBSERVED_NEUTRAL (src/backtest/executor.py, on a branch, for the next design rerun) so that the averages of an unobserved series are exactly neutral.
Problem (updated 2026-10-22, checked on the 13-bar design): only the volatility-demand series stop early, and the text discloses none of the tails.
VIX, VIX3M, VVIX: the vendor file ends on 12 February 2024 (data/vix_and_voldemand.parquet); the design splices Yahoo's hourly closes from 13 February to 30 April 2024, 54 trade days (src/data/loading.py, SPLICE_VIX_FREE_FEED), each hourly close carried into the next half-hour bar. Agreement where both exist: on the hourly overlap (December 2023 to February 2024) the Yahoo close equals the vendor print on 96% (VIX), 88% (VIX3M) and 100% (VVIX) of stamps (results/data_audit_2026-10-15/vix_yahoo_overlap.csv); at the daily close, 2006–2023, Yahoo's close equals the vendor bar holding the 16:15 close on 96% of days for VIX and 93% for VIX3M (99% and 98% within 0.1%; daily log changes correlate 0.9998 and 0.995; results/data_audit_2026-10-22/yahoo_vs_vendor_daily.csv; VVIX has no vendor bar after 16:00). In the design these series are live to the end of the span (availability about 0.85–0.90). *FOMC:* the release file stops at the 1 November 2023 statement; the three scheduled statements to 30 April 2024 are added (FOMC_CALENDAR_FIX), so no FOMC row is unobserved. *Volatility demand:* the last print is on 31 August 2023. The forward fill is capped at 26 bars (APPLY_FFILL_LIMIT, since 15 October 2026); after that the series is unobserved, its availability averages $\bar{a}^{(k)}$ fall to zero, and its value averages hold the neutral fill, which the rolling scaler keeps rescaling as its window moves (a slow drift, e.g. from about -0.28 to -0.16 for the 1-interval average, September 2023 to April 2024). UNOBSERVED_NEUTRAL sets them to exactly 0 there. Disclosure for the text: the end of the volatility-demand feed and the cap; the VIX-family splice and its agreement.

– *[name the vendor of each file].*

E2. Decision: name the vendor, the price field, the roll, and the bin edges.
Problem: the vendor, the price inside the one-minute log change, the day whose volume picks the front month, and the clock edges of every bin except 15:30–16:00 and the weekend break are missing, so the panel cannot be rebuilt.
The futures moments are now sums of one-minute changes in the log price of the front-month contract, the most active. The text does not say whether that price is a last trade, a midpoint, or a settlement, and it does not name the vendor of the file. It does not say which day's volume makes a contract the most active. The closing bin is 15:30–16:00 ET, and the grid drops Friday 20:00 through Sunday 18:30. Each other bin is a half hour ending on a half-hour stamp. Those stamps are not listed, so the $M = 48$ bins cannot be rebuilt from the text.

- **Intraday grid and data cleaning.**
  – We construct the panel on a 30-minute grid starting from the first observed realized variance. An interval i on day t is retained in the interval grid $\mathcal{G}$ if and only if its realized variance is observed and it falls outside the weekend trading break:

$$(t,i) \in \mathcal{G} \iff \mathrm{RV}_{t,i} \text{ observed and } (t,i) \notin (\text{Fri } 20{:}00,\ \text{Sun } 18{:}30).$$

  – Lag operators step backward sequentially through the retained intervals in $\mathcal{G}$ in chronological order, spanning overnight sessions and trading-day boundaries; thus, $(t, i-j)$

refers to the j-th retained interval preceding (t, i), which may lie on an earlier trading day.

- Exchange holidays and early-close sessions are retained in the futures panel, whereas early-close days are excluded from the option-trading evaluation sample.
- For each exogenous series u, we first record a binary availability (non-missingness) indicator prior to any imputation:

$$a_{t,i} = \mathbf{1}\{u \text{ observed in interval } (t, i)\} \in \{0, 1\}.$$

- Subsequently, $u_{t,i}$ is forward-filled (carrying the last observed value forward without a staleness cap) for the purpose of computing trailing baseline statistics, after which every interval with $a_{t,i} = 0$ is replaced by the causal median imputation described in step 4 of Section 4.
- The diurnal adjustment factor $s^2_{t,i}$ and the winsorized target $\tilde{y}_{t,i}$ are formally defined in Section 5.1, with the corresponding rolling window lengths and quantile thresholds reported in Appendix B.1.

- **Estimation and evaluation samples.** Appendix Table 4 lists the samples. The estimation samples and the evaluation span of the forecasts are set out with the experiments (Section 6.1), the option-trading sample with the application (Section 7.1), and the closing-interval sample with the closing-interval models (Section G.5).

4 Predictors

- **Eight groups of exogenous predictors (46 series in total).** As listed in Appendix Table 6, the exogenous series span eight economic categories: realized moments (7), liquidity (14), equal-weighted market moments (6), value-weighted market moments (6), sentiment (3), implied volatility (3), volatility demand (4), and FOMC announcements (3). Section 3 describes the construction of each category, and Appendix Table 6 provides the definition of each individual series:

$$46 = 7 + 14 + 6 + 6 + 3 + 3 + 4 + 3.$$

- Specifically, the 3 FOMC series are defined as

$$\mathbf{1}\{\text{announcement on day } t\}, \quad 1/(1 + n^{+}_{t,i}), \quad 1/(1 + n^{-}_{t,i}),$$

- where $n^{\pm}_{t,i}$ denotes the number of intervals until the next $(+)$ and since the last $(-)$ scheduled announcement.

- **Predictors constructed from each series.**
  - Following Corsi (2009), we construct HAR-type trailing averages of each transformed series $\tilde{u}$ over horizons of k samples (one sample is one 30-minute grid interval), lagged by one sample:

$$\bar{u}^{(k)}_{t,i} = \frac{1}{k} \sum_{j=1}^{k} \tilde{u}_{t,i-j}, \qquad k \in \mathcal{K},$$

  where the horizon set is

$$\mathcal{K} = \{1, 5, 25, 125, 625, 3125\} \quad \text{(base 5, 6 horizons)}.$$

- Here, k counts samples, the retained intervals of the full intraday grid $\mathcal{G}$ (including overnight sessions), in the intraday and the closing-interval model alike. The closing-interval specification uses the same horizon set (Appendix G.2).
- In addition, we compute the corresponding k-interval trailing averages of the availability indicator, denoted $\bar{a}_{t,i}^{(k)}$.
- The intraday model does not use trailing averages of the activity indicator $\chi_{t,i}$ (step 5 of the transformations below). The closing-interval model does: $\bar{\chi}_{t,i}^{(k)}$, for its mode-inflated series (Appendix G.2).
- For notational simplicity in our tables and figures, we suppress the overbar and the interval index i ($i = c$ for the scored forecast), writing $u_t^{(k)}$ for series u averaged over k samples; for instance, $\texttt{vix}_t^{(25)}$ denotes the transformed VIX series averaged over the 25 samples ending at 15:30 on day t. We refer to each exogenous series by its column name in the data (Table 6): the implied volatility series are `vix`, `vix3m`, and `vvix`; the suffixes `_ewstock` and `_vwstock` mark the equal- and value-weighted constituent aggregates (e.g., `ofi_ewstock`, `sumret2_vwstock`, and `spread_vwstock`); and the announcement predictors are `fomc_day`, `fomc_until_inv`, and `fomc_since_inv`, while writing $\bar{a}_t^{(k)}(u)$ and $\bar{\chi}_t^{(k)}(u)$ for the trailing averages of the availability and activity indicators.

- **HAR benchmark.** The HAR benchmark predictor set consists of the HAR components of the target, $\bar{y}_{t,i}^{(k)}$—constructed via the same trailing-average operator as $\bar{u}_{t,i}^{(k)}$ applied to the winsorized target (6 columns in the intraday model)—together with the calendar vector $\boldsymbol{D}_{t,i}$, which comprises five day-of-the-week dummy variables, the hour of the interval-end timestamp, and the overnight session indicator.

E9. Decision: the calendar stops at the overnight indicator, and that indicator stays only if the overnight intervals stay in the training sample (note E15).
Problem: the columns after the overnight indicator are out of the specification, and the printed column counts and Table 1 still include them. Whether the overnight indicator itself stays depends on the sample in note E15.
The sentence above keeps three calendar pieces: the five weekday dummies, the hour of the bar-end stamp, and the overnight indicator. Every column that used to sit after the overnight indicator is out of the specification. That is the opening-interval indicator, the closing-interval indicator, expiration, expiration week, quadruple witching, the quadruple-witching close, month end, quarter end, days to expiration, and the products of the HAR components with the opening and closing indicators. The overnight indicator's own place depends on which sample note E15 keeps. On all 48 intervals it equals 1 off the cash session and again on the 16:00 bar, so it varies across the window and a regression can estimate it. On the 13 regular-hours bars it equals 1 only on the 16:00 bar, so inside that sample it is constant on every bar but the close. The column counts in the formulas below, and the fits in Table 1, were built with the dropped columns still in the matrix.

- **Sequential feature transformations.** Each exogenous series undergoes the following seven transformations in sequence:

1. *Diurnal adjustment:* we divide each raw series $u_{t,i}$ by a time-of-day scale factor $b_{t,i}^{(u)}$ estimated over the same prior-day set $\mathcal{D}_{t,i}$ as $s_{t,i}^2$ and bounded relative to its own trailing median:

$$\tilde{u}_{t,i} = \frac{u_{t,i}}{b_{t,i}^{(u)}}, \qquad b_{t,i}^{(u)} = \begin{cases} \mathrm{clip}\big(\mathrm{mean}_{t' \in \mathcal{D}_{t,i}}\, u_{t',i},\ 0.5\,\rho_{t,i},\ 2\rho_{t,i}\big), & u \geq 0, \\ \max\big\{\mathrm{clip}\big(\mathrm{sd}_{t' \in \mathcal{D}_{t,i}}\, u_{t',i},\ 0.5\,\rho_{t,i},\ 2\rho_{t,i}\big),\ 0.1\,\bar{\sigma}_{t,i}^{(u)}\big\}, & u \text{ signed}, \end{cases}$$

   - where $\rho_{t,i}$ denotes the median of the 250 most recent unbounded diurnal factors for interval i (requiring at least 20 observations),

- $\bar{\sigma}_{t,i}^{(u)}$ denotes the median of the nonzero trailing standard deviations $\mathrm{sd}_{t'' \in \mathcal{D}_{t',i'}} u_{t'',i'}$ over all intervals (t', i') up to (t, i), every intraday interval pooled (before the first such standard deviation the adjusted value is undefined and is imputed as in step 4),
- a series is signed when the quantity it measures can be negative, not when its sample has a negative value (Table 6),
- $b_{t,i}^{(u)}$ is set to 1 whenever the factor of a nonnegative series is zero or undefined, and
- no diurnal adjustment is applied to the implied volatility indices (VIX, VIX3M, and VVIX), the StockTwits sentiment series, or the FOMC series.
- The closing-interval forecasts predate this floor and this sign rule: they take $\bar{\sigma}^{(u)}$ as the full-sample median of the same standard deviations, and call a series signed when any value in its full sample is negative, which signs the value-weighted quoted spread through a single invalid print (Table 6).

2. *Variance-stabilizing transformation:* we apply a domain-specific mapping $\tilde{u}_{t,i} \leftarrow g(\tilde{u}_{t,i})$ according to the first matching rule in

$$g(v) = \begin{cases} \sqrt{v} & \text{squared returns, realized variance, turnover,} \\ & \text{bipower variation, effective spread,} \\ \operatorname{sign}(v)\,|v|^{1/2} & \text{autocovariance,} \\ v^{1/3} & \text{cubed returns,} \\ \operatorname{sign}(v)\,|v|^{1/4} & \text{fourth-power returns,} \\ v & \text{other signed series, absolute returns,} \\ \log \max(v, 10^{-12}) & \text{otherwise.} \end{cases}$$

   - A signed series without a specific power-root rule is routed at interval (t, i) to an inverse hyperbolic sine transformation of the raw series when its standard-deviation divisor sat at its floor on at least 50% of its intervals since its first print and before (t, i), at least 250 of them. This routes the volatility-demand series, on almost all of their intervals (Table 6), and no other series:

$$\tilde{u}_{t,i} = \operatorname{asinh}(u_{t,i}/\kappa_{t,i}), \qquad \kappa_{t,i} = \max\{e_{t,i-1},\ 0.01 \operatorname{med}_{j \geq 1} e_{t,i-j}\},$$

   - where $e_{t,i}$ denotes an exponentially weighted moving average of $|u|$ computed over nonzero intervals ($u \neq 0$) with a half-life of 12,000 such intervals (nominally 250 days of 48 intervals).
   - The FOMC statement-day dummy passes through unchanged, with no transformation and no winsorization (step 3).
   - The closing-interval forecasts predate these rules: they route a series by the share over its full sample, the intervals before its first print included, which routes the same volatility-demand series on every interval; they fill $\kappa_{t,i}$ with its full-sample median where it is not yet defined; and they take the logarithm of the FOMC statement-day dummy and winsorize it.

3. *Winsorization:* each series is clipped at the same rolling quantile thresholds used for the target (see Section 5.1), evaluated over the series' own preceding intervals.
4. *Missing-value imputation:* any missing observation is imputed using the expanding median of the series' previously observed transformed values, while retaining the availability

indicator $a_{t,i}$ as a separate predictor:

$$\tilde{u}_{t,i} \leftarrow \begin{cases} \tilde{u}_{t,i}, & a_{t,i} = 1, \\ \operatorname{med} \mathcal{A}_{t,i}, & a_{t,i} = 0,\ \mathcal{A}_{t,i} \neq \emptyset, \\ \tilde{u}_{(1)}, & a_{t,i} = 0,\ \mathcal{A}_{t,i} = \emptyset, \end{cases}$$

- where $\mathcal{A}_{t,i} = \{\tilde{u}_{t',i'} : (t',i') \prec (t,i),\ a_{t',i'} = 1\}$, with $(t',i') \prec (t,i)$ indicating that (t',i') is a retained interval strictly prior to (t,i), and
- $\tilde{u}_{(1)}$ denotes the transformed value on the series' first observed interval.

5. *Activity indicators:* a series is mode-inflated at interval (t,i) when it has at least 250 observed prints strictly before (t,i) and their most frequent value, $m_{t,i}$, accounts for at least 20% of them. For each series that is mode-inflated at some interval, we define a binary activity indicator from its observed prints: on every interval at which the series is mode-inflated,

$$\chi_{t,i} = \begin{cases} \mathbf{1}\{u_{t,i} \neq m_{t,i}\}, & a_{t,i} = 1, \\ \chi^{\circ}_{t,i}, & a_{t,i} = 0, \end{cases}$$

where $\chi^{\circ}_{t,i}$ is the more frequent value of χ on the series' earlier observed intervals (0 on a tie), and $\chi_{t,i} = 0$ on the other intervals.

- The mode-inflated class and the mode are decided from earlier prints only, and the indicator is read from observed prints, not from the forward-filled series. Under this rule, our specification, 5 series are mode-inflated at some interval: the number of observations, the 3 StockTwits series (on part of the sample) and the FOMC statement-day dummy; the volatility-demand series never are.
- The pooled intraday models of Table 1 carry no activity averages $\bar{\chi}^{(k)}$; the indicator enters the closing-interval models (Appendix G.2).
- The closing-interval forecasts we report predate this rule and use a full-sample one: a series is mode-inflated when its full-sample mode accounts for at least 20% of all intervals, and $\chi_{t,i} = \mathbf{1}\{u_{t,i} \neq \text{mode}(u)\}$ is read on the forward-filled series. Under that rule the 6 series marked in Table 6 carry the indicator (the number of observations, the 4 volatility-demand series and the FOMC statement-day dummy) and, in the fitted forecasts, so does the FOMC statement-interval dummy, since dropped from the predictors.

6. *HAR aggregation:* we form the trailing averages $\bar{u}^{(k)}_{t,i}$, $\bar{a}^{(k)}_{t,i}$, $\bar{\chi}^{(k)}_{t,i}$ (closing-interval model only), and $\bar{y}^{(k)}_{t,i}$ across horizons $k \in \mathcal{K}$.
7. *Robust standardization:* under the pooled intraday scheme, we standardize the trailing averages of series values over the rolling estimation window using the median and interquartile range (with lower floors). Specifically, the median and interquartile range are computed exclusively over intervals with $a_{t',i'} = 1$ inside the rolling window $\mathcal{W}_{t,i}$ (see Appendix B.2), re-evaluated every 480 intervals and held constant until the next grid point. Let $M^{(k)}_{t,i}$ denote the maximum such interquartile range observed across all grid points up to (t,i). We apply a 0.01 relative floor in the denominator (note that $\text{IQR}^{*(k)}$, defined in Appendix G.4, does not enter this pooled denominator):

$$z^{(k)}_{t,i} = \begin{cases} \dfrac{\bar{u}^{(k)}_{t,i} - \operatorname{med}\{\bar{u}^{(k)}_{t',i'} : (t',i') \in \mathcal{W}_{t,i},\ a_{t',i'} = 1\}}{\max\{\text{IQR}^{(k)}_{t,i},\ 0.01\, M^{(k)}_{t,i}\}} & \text{if at least 512 such intervals,} \\ 0 & \text{otherwise,} \end{cases}$$

$$\mathrm{IQR}_{t,i}^{(k)} = \mathrm{IQR}\{\bar{u}_{t',i'}^{(k)} : (t',i') \in \mathcal{W}_{t,i},\ a_{t',i'} = 1\}.$$

- On the initial estimation window, both statistics are computed using that window's own rows; an interquartile range below 10^{-12} is replaced by 1 on the initial window and, at subsequent grid points, retains the previous scale factor.
- The availability averages $\bar{a}^{(k)}$ are left unstandardized.
- The target HAR components and calendar variables do not use the availability-restricted denominator above; instead, their median and interquartile range are computed across all rows of $\mathcal{W}_{t,i}$ with $\mathrm{IQR}^{*(k)} = 0$ ($\mathrm{IQR}^{*(k)}$ is defined in Appendix G.4), subject to the same 0.01 floor relative to the running maximum interquartile range of that column.
- For a series whose first print comes after the first day of the sample, $z_{t,i}^{(k)} = 0$ until every interval of the window $\mathcal{W}_{t,i}$ is live, an interval being live when its k-interval average holds no interval before the day of the series' first print.
- $z_{t,i}^{(k)} = 0$ whenever the k-interval average holds no observed interval, $\bar{a}_{t,i}^{(k)} = 0$, so that the averages of a feed that has stopped are exactly neutral.
- This step equalizes each column's interquartile range, not its variance. The ratio of a column's standard deviation to its interquartile range varies with its tails, its skewness and its share of zeros; the availability averages $\bar{a}^{(k)}$ keep their $[0,1]$ scale; and the scale factors come from the rolling grid rather than from the sample on which a model is estimated. The step serves the data: it limits the influence of outliers and keeps each series comparable over time. The penalized linear models therefore also standardize every column within their own estimation window (Section 5.2), so that the penalty treats every column alike.

S4d. Decision: the late-start and no-observation rules of this step are the specification (src/backtest/executor.py LATE_START_WARMUP, merged 21570f0, and UNOBSERVED_NEUTRAL, merged 5403acd; both 22 October 2026; note E28). A rerun of the pooled intraday models under them is approved and waits for the window choice.
Problem: the printed pooled forecasts of Table 1 predate both rules (their designs were built on 20 October 2026). There a series that starts late takes the scaler's output from its first window on, and the value averages of a feed that has stopped (volatility demand after 31 August 2023) hold the rescaled neutral fill instead of 0.

- **Column screening.** At each estimation origin, we screen out constant and duplicate predictors over the active estimation window $\mathcal{W}_t$ by retaining the index set

$$\mathcal{J}_t = \{j :\ \boldsymbol{x}_{\cdot j} \text{ not constant},\ \boldsymbol{x}_{\cdot j} \neq \boldsymbol{x}_{\cdot j'}\ \forall\, j' < j,\ j' \in \mathcal{J}_t\}, \qquad \boldsymbol{x}_{\cdot j} = (x_{t'j})_{t' \in \mathcal{W}_t}.$$

  - Specifically, column equality is evaluated bitwise, and whenever two columns are identical on $\mathcal{W}_t$, the first column in order is retained.

- **Predictor set dimensions.** The column counts of the closing-interval specification are given in Appendix G.2.

  - In the pooled intraday specification, the full predictor set comprises 586 candidate columns, constructed from 46 exogenous series, 6 horizons, and 12 interaction terms between the target HAR components and the opening- and closing-interval dummy variables. Note that these column counts still reflect the calendar columns and interaction terms discussed in note E9 above:

$$586 = \underbrace{6 + 16}_{\bar{y}^{(k)},\ \boldsymbol{D}} + \underbrace{12}_{\bar{y}^{(k)} \times \text{open, close}} + \underbrace{46 \times 6}_{\bar{u}^{(k)}} + \underbrace{46 \times 6}_{\bar{a}^{(k)}}.$$

Here, the 46 intraday exogenous series encompass every category in Section 3, the FOMC group (3 series) included. Appendix Table 6 lists each individual series.

E35. Decision: base 5 across the paper, and the pooled model includes the FOMC series (user, 2026-10-19).
Problem: 586 and the formula above now count the base-5 pooled panel with the FOMC series, built on 19 October 2026, less its activity averages (note E6) and before the cut of note E9; no printed forecast was fitted on that design.
The pooled trees printed before 22 October 2026 and the intraday OLS were fitted on the base-2 panel without the FOMC series, 1,144 columns from 12 horizons and 43 series; the pooled trees are now refit on the 576-column design on the 13 regular-hours bars described at the end of this note, the design of the pooled linear models of Table 1 (note E34). The base-5 panel (experiments/pooled_trees_base5.py, design stage: the repo's own preparation with the data fixes of 15 October 2026, the FOMC calendar fix and the causal mode rule) has 6 horizons, 46 series and 12 interaction columns. Under the causal rule 5 series carry an activity indicator, the number of observations, the three StockTwits series and the FOMC statement-day dummy, where the old full-sample rule gave 5 on the panel without FOMC, the number of observations and the four volatility-demand series (note S4a). With the activity averages dropped (E6) and the calendar cut after the overnight indicator together with the interaction columns (E9), the design has 577 columns on all 48 intervals and 576 on the 13 regular-hours bars, where the overnight indicator is dropped too. Those two counts have no macro yet.
Since 22 October 2026 (third-list item 5) the FOMC statement-interval dummy is out of the predictor set (src/data/loading.py, EXCLUDE_FOMC_RELEASE): 46, 586, 5, the FOMC group, Table 6 and the design table (generated/rehaul_dm_design.tex, not in the paper) count the panel without its columns. The 577- and 576-column designs above, and so every printed pooled forecast, still contain its averages and availability averages; on the 13-bar design the dummy is a live 0/1 column (one on the statement bar of each scheduled meeting). A rerun on the design without it is needed before the counts and the printed forecasts agree.

- **Causal construction and absence of look-ahead.** Every rolling statistic indexed by (t, i) is computed strictly from intervals prior to (t, i), ensuring that every predictor of interval (t, i) is in the information set at its forecast origin (15:30 ET for the scored closing interval $i = c$). This includes the HAR trailing averages, the diurnal factor $s^2_{t,i}$, the winsorization quantiles, the median baseline $\rho_{t,i}$, the standardization moments, and the active column set $\mathcal{J}_t$.
  - In the pooled intraday forecasts, the standard-deviation floor $\bar{\sigma}^{(u)}_{t,i}$ and the inverse hyperbolic sine routing are computed from earlier intervals, and the sign of each series is read from what it measures, not from the full sample (steps 1 and 2); the activity indicator does not enter them (step 5).
  - We disclose the full-sample quantities used in preprocessing the closing-interval forecasts, which predate those rules: the mode that defines $\chi_{t,i}$ under the rule those forecasts were built with (step 5), the standard-deviation floor $\bar{\sigma}^{(u)}$ and the sign of each series (step 1), and the rule assigning series to the inverse hyperbolic sine transformation (step 2).
  - As noted in Section 3, the FOMC proximity variables rely solely on the pre-announced meeting schedule, which is publicly known ahead of each interval.
  - Any missing interval prior to a series' initial observation is filled with that first observed print (step 4); under the late-start rule of step 7 the standardized averages that hold such intervals are set to 0.
  - Finally, standardization statistics on the first estimation window are computed exclusively from burn-in observations.

5 Methods

5.1 Problem formulation

- **Rolling-window out-of-sample exercise.** We conduct a rolling-window out-of-sample forecasting exercise in which each model is re-estimated every period using only predictors in the information set $\mathcal{F}_t$ available at the 15:30 ET forecast origin. A period corresponds to a 30-minute intraday interval under our pooled intraday scheme (Appendix B.2).
- **Notation.** We index trading days by $t = 1, \ldots, T$ and the 30-minute intraday intervals of day t (including the overnight session) by $i = 1, \ldots, M$, where $i = c$ denotes the closing half-hour interval (15:30–16:00 ET).
  - Realized variance $\mathrm{RV}_{t,i}$ and the lag operator $(t, i-j)$ are defined in Section 3. Our primary forecast target is the realized variance of the closing interval, $\mathrm{RV}_t \equiv \mathrm{RV}_{t,c}$.
  - We denote by $\mathcal{F}_t$ the information set containing all data observed through 15:30 ET on day t, and write $\mathbb{E}_t[\cdot] \equiv \mathbb{E}[\cdot \mid \mathcal{F}_t]$ and $\mathrm{Var}_t(\cdot) \equiv \mathrm{Var}(\cdot \mid \mathcal{F}_t)$ for the corresponding conditional expectation and variance.
- **Diurnal adjustment.** Following Andersen and Bollerslev (1997), we remove the intraday diurnal pattern using the trailing sample mean of realized variance in the same time-of-day interval strictly prior to day t:

$$s^2_{t,i} = \frac{1}{|\mathcal{D}_{t,i}|} \sum_{t' \in \mathcal{D}_{t,i}} \mathrm{RV}_{t',i},$$

$$\mathcal{D}_{t,i} = \{\text{last } n_{\mathrm{d}} \text{ trading days } t' < t \text{ with interval } i\}, \qquad |\mathcal{D}_{t,i}| \geq \underline{n}_{\mathrm{d}}.$$

  - For intervals with $|\mathcal{D}_{t,i}| < \underline{n}_{\mathrm{d}}$, we set $s^2_{t,i} = 1$ and exclude the interval from evaluation. The window parameters n_{d} and $\underline{n}_{\mathrm{d}}$ are given in Appendix B.1.
- **Regression target.** To stabilize the variance, we take the square root of the diurnally adjusted realized variance and winsorize the resulting series at rolling lower and upper empirical quantiles:

$$y_{t,i} = \sqrt{\mathrm{RV}_{t,i}/s^2_{t,i}}, \qquad \mathcal{Y}_{t,i} = \{y_{t,i-j} : j = 1, \ldots, n_{\mathrm{w}}\},$$

$$\tilde{y}_{t,i} = \mathrm{clip}(y_{t,i},\ \underline{Q}_{t,i},\ \overline{Q}_{t,i}).$$

  - Here $\underline{Q}_{t,i}$ and $\overline{Q}_{t,i}$ denote the lower and upper empirical quantiles computed over $\mathcal{Y}_{t,i}$, which pools the preceding n_{w} intervals across all times of day, and clip truncates its first argument to that range. The quantile levels and window length n_{w} are reported in Appendix B.1.
  - All models are estimated on the winsorized target $\tilde{y}_{t,i}$, whereas out-of-sample evaluation is conducted against the unwinsorized realized variance $\mathrm{RV}_{t,c}$.
  - For the scored interval $i = c$ (the closing half hour that every model forecasts), we write $y_t \equiv y_{t,c}$ and $\tilde{y}_t \equiv \tilde{y}_{t,c}$.
- **Back-transformation.** To obtain a variance forecast in raw units, we square the fitted volatility and multiply by the diurnal factor without post-estimation recalibration:

$$\hat{\sigma}^2_t = \hat{y}_t^{\,2}\, s^2_{t,c},$$

  where $\hat{y}_t \equiv \hat{y}_{t,c}$ is the forecast of the winsorized target $\tilde{y}_t$, not of the unclipped y_t, formed using $\mathcal{F}_t$.

– Because $\mathbb{E}_t[y_t^2] - (\mathbb{E}_t[y_t])^2 = \mathrm{Var}_t(y_t) \geq 0$, squaring the level forecast introduces a downward Jensen inequality bias that we leave uncorrected (note E33).

- **Rolling estimation window.** Under the intraday scheme, the estimation window at origin (t, i) is $\mathcal{W}_{t,i}$ (Appendix B.2), consisting of the W most recent intervals prior to (t, i). We let $\mathcal{I}$ denote the subset of intraday intervals included in estimation.
- **Hyperparameter-selection splits.** In line with Gu et al. (2020), we partition the rolling window forward in time into an estimation subsample, an embargo gap of n_{gap} observations, and a validation subsample of n_{val} observations:

$$\begin{aligned}\mathcal{W}_{t,i}^{\text{est}} &= \{(t,\, i-W), \ldots, (t,\, i-n_{\text{val}}-n_{\text{gap}}-1)\},\\ \mathcal{W}_{t,i}^{\text{gap}} &= \{(t,\, i-n_{\text{val}}-n_{\text{gap}}), \ldots, (t,\, i-n_{\text{val}}-1)\},\\ \mathcal{W}_{t,i}^{\text{val}} &= \{(t,\, i-n_{\text{val}}), \ldots, (t,\, i-1)\}.\end{aligned}$$

  Target interval (t, i) never enters any of these sets. The lengths n_{gap} and n_{val} are given in Appendix B.1.
- **Estimation objective.** Each model is estimated by minimizing the sum of squared errors on the scale of $\tilde{y}$. At forecast origin (t, i), the parameter vector $\boldsymbol{\theta}$ is fitted over the full rolling window $\mathcal{W}_{t,i}$:

$$\hat{\boldsymbol{\theta}}_{t,i} = \arg\min_{\boldsymbol{\theta}} \sum_{(t',i')\in\mathcal{W}_{t,i}} \left(\tilde{y}_{t',i'} - \hat{y}_{t',i'}(\boldsymbol{\theta})\right)^2.$$

  – Here (t', i') indexes the historical intervals in the window $\mathcal{W}_{t,i}$, and $\boldsymbol{\theta}$ collects the model parameters.
  – During hyperparameter tuning only, candidate configurations are estimated on $\mathcal{W}_{t,i}^{\text{est}}$ rather than $\mathcal{W}_{t,i}$ and evaluated on $\mathcal{W}_{t,i}^{\text{val}}$.

5.2 Models

- **HAR and the HAR benchmark via OLS.** Following Corsi (2009), we specify the linear heterogeneous autoregressive model for each intraday interval (t, i) as

$$\hat{y}_{t,i} = \beta_0 + \sum_{k\in\mathcal{K}} \beta^{(k)}\, \bar{y}_{t,i}^{(k)} + \boldsymbol{\delta}^\top \boldsymbol{D}_{t,i}.$$

  – The horizon set $\mathcal{K}$ and the lagged volatility averages $\bar{y}_{t,i}^{(k)}$ are defined in Section 4, where $\bar{y}_{t,i}^{(k)}$ is the mean of the winsorized target over the k half-hour intervals preceding (t, i), including overnight intervals. The scored forecast is the closing-interval prediction $\hat{y}_t \equiv \hat{y}_{t,c}$.
  – The vector $\boldsymbol{D}_{t,i}$ collects the calendar dummy variables with coefficient vector $\boldsymbol{\delta}$; a pure HAR specification restricts $\boldsymbol{\delta} = \mathbf{0}$. Our HAR benchmark ($\mathcal{X} = \text{HAR}$) includes both the HAR components and the calendar dummies.
  – The full predictor vector $\boldsymbol{x}_{t,i} \in \mathbb{R}^p$ ($\mathcal{X} = \text{Full}$) augments the HAR benchmark with the multi-horizon averages $\bar{u}_{t,i}^{(k)}$ ($k \in \mathcal{K}$) and the corresponding availability and activity indicators for every exogenous series u (Section 4).
- **Penalized linear regressions: ridge (Hoerl and Kennard, 1970), lasso (Tibshirani, 1996), and elastic net (Zou and Hastie, 2005).** To accommodate the high-dimensional predictor vector $\boldsymbol{x}_{t,i}$, we augment the residual sum of squares with an ℓ_1, ℓ_2, or combined

penalty, following the penalized HAR formulation of Zhang et al. (2024), and update the coefficient estimates as the estimation window rolls forward:

$$\hat{\boldsymbol{\beta}}^{\text{ridge}}_{t,i} = \arg\min_{\beta_0,\boldsymbol{\beta}} \sum_{(t',i')\in\mathcal{W}_{t,i}} (\tilde{y}_{t',i'} - \beta_0 - \boldsymbol{x}^{\top}_{t',i'}\boldsymbol{\beta})^2 + \lambda \sum_{j=1}^{p} \omega_j \, \beta_j^2,$$

$$\hat{\boldsymbol{\beta}}^{\text{enet}}_{t,i} = \arg\min_{\beta_0,\boldsymbol{\beta}} \sum_{(t',i')\in\mathcal{W}_{t,i}} (\tilde{y}_{t',i'} - \beta_0 - \boldsymbol{x}^{\top}_{t',i'}\boldsymbol{\beta})^2 + \lambda\alpha \sum_{j=1}^{p} \omega_j |\beta_j| + \frac{\lambda(1-\alpha)}{2} \sum_{j=1}^{p} \omega_j \, \beta_j^2,$$

where $\lambda > 0$ governs overall penalty strength, $\alpha \in [0,1]$ is the elastic-net mixing parameter, and $\omega_j \geq 0$ is the predictor-specific penalty weight (with the intercept β_0 unpenalized, $\omega_0 = 0$). For ridge, λ coincides directly with the grid value reported in Table 5. For lasso and elastic net, Table 5 reports the `scikit-learn` penalty parameter, so the penalty multiplier λ in the unnormalized objective above equals $2n$ times the tabulated value, where n is the number of observations in the window.

- **Differential block penalty.** When estimating ridge, lasso, and elastic net on the full predictor set, we penalize the HAR benchmark predictors 100 times more lightly than the exogenous predictors:

$$\omega_j = \begin{cases} \underline{\omega} = 0.01 & j \in \text{HAR benchmark}, \\ 1 & j \text{ an additional predictor}, \\ 0 & j = 0 \text{ (intercept)}. \end{cases}$$

- When a penalized regression is fitted on the HAR benchmark alone, we apply a uniform penalty weight $\omega_j = 1$ across all predictors.
- **Standardization within the estimation window.** The penalties apply to standardized coefficients (Tibshirani, 1996; Friedman et al., 2010). Each predictor column is centered and scaled by its mean m_j and standard deviation s_j over $\mathcal{W}_{t,i}$; the objectives above are solved for the coefficients γ_j of the standardized columns $(x_j - m_j)/s_j$, with the weights ω_j, and transformed back, $\beta_j = \gamma_j/s_j$, the intercept absorbing the centering. Without this step the penalty on a column's contribution to the forecast would scale with $1/s_j^2$ for ridge and $1/s_j$ for the lasso, so columns with little spread, such as long-horizon averages, would be penalized far more heavily for no economic reason. The robust standardization of Section 4 does not achieve this: it equalizes interquartile ranges rather than variances, and its scale factors come from the rolling grid rather than from the estimation window. For ridge, m_j and s_j are updated at every step; for the lasso and elastic net, at every cold restart of the homotopy described below, and held fixed in between.
- A column of an exogenous series enters a fit only if the series' own one-interval column differs from its modal value on at least 250 intervals of $\mathcal{W}_{t,i}$, the evidence requirement of the activity indicator (Section 4). A k-interval average repeats each observation k times, so its own count would overstate the evidence. Columns that fail are screened out like constant columns (below).
- An availability average over $k > 1$ intervals is a proportion, with sampling resolution $\sqrt{p(1-p)/k}$ at its window mean p. It is screened out of a fit when its standard deviation over $\mathcal{W}_{t,i}$ does not exceed that resolution, or when the standard deviation of its difference from an earlier retained availability average of the same horizon does not exceed $\sqrt{2}$ times the resolution at the pair's mean: its movement is then no larger than independent

missing observations would produce, so it carries no detectable information. These are mostly the long-horizon availability averages of regular-hours series, which record the trading calendar rather than the series.

– The solvers work directly on the standardized columns, which gives the same estimates as penalizing the raw columns with weights $\omega_j s_j$ and $\omega_j s_j^2$ but keeps the systems they solve well conditioned.

STD. Decision: standardization within the estimation window, the availability screen and the standardized-coordinate solvers are the specification (STANDARDIZE_PENALTY, merged 2430543 and 8898522, 22 October 2026).
Problem: the printed pooled forecasts predate all three; the rerun is pending (notes to the authors, "Rerun pending").

– The lasso corresponds to $\alpha = 1$, and our elastic net specification fixes $\alpha = 0.5$.

– For ridge regression, we update $(\boldsymbol{X}_{t,i}^\top \boldsymbol{X}_{t,i} + \lambda \operatorname{diag}(\boldsymbol{\omega}))^{-1}$ via exact rank-one downdates and updates as observations exit and enter $\mathcal{W}_{t,i}$, where $\boldsymbol{X}_{t,i}$ stacks the rows $(1, \boldsymbol{x}_{t',i'}^\top)$ for $(t', i') \in \mathcal{W}_{t,i}$ and $\boldsymbol{\omega} = (\omega_0, \ldots, \omega_p)^\top$.

– For the lasso and elastic net, we update the active set and coefficient vector via exact online homotopy (Garrigues and El Ghaoui, 2008), adding the incoming observation and dropping the outgoing observation at each step.

– At every estimation step, only the screened predictor subset $\mathcal{J}_{t,i}$ enters the solver: columns that are constant within $\mathcal{W}_{t,i}$ or exact duplicates of an earlier retained column are omitted. In models estimated solely on the closing interval, this screen automatically removes the hour-of-day and overnight indicators (Section 4).

- **Tree ensembles.** To capture potential nonlinearities and high-order interactions among predictors, we consider two gradient-boosted tree implementations (LightGBM and XGBoost) and random forests.

  – **Gradient boosting** (Friedman, 2001), implemented via LightGBM (Ke et al., 2017) and XGBoost (Chen and Guestrin, 2016), minimizes the squared-error loss $\frac{1}{2}(\tilde{y}-F)^2$ over boosting rounds $b = 1, \ldots, B$ with learning rate η. Each regression tree f_b partitions the predictor space into J leaves with leaf weights $\boldsymbol{w}$, where leaf j contains the observation subset $\mathcal{R}_j$ of size $n_j = |\mathcal{R}_j|$. When the ℓ_1 leaf penalty $\lambda_1^{\text{tree}} = 0$, the additive update and optimal leaf weight take the form

$$F_b = F_{b-1} + \eta f_b,$$
$$w_j^* = -\frac{\sum_{(t',i') \in \mathcal{R}_j} (F_{b-1}(\boldsymbol{x}_{t',i'}) - \tilde{y}_{t',i'})}{n_j + \lambda_2^{\text{tree}}}.$$

  At round b, tree f_b is grown to minimize the regularized second-order objective

$$\sum_{(t',i') \in \mathcal{W}_{t,i}} \left[(F_{b-1}(\boldsymbol{x}_{t',i'}) - \tilde{y}_{t',i'}) f_b(\boldsymbol{x}_{t',i'}) + \tfrac{1}{2} f_b(\boldsymbol{x}_{t',i'})^2 \right] + \gamma^{\text{tree}} J + \lambda_1^{\text{tree}} \|\boldsymbol{w}\|_1 + \tfrac{1}{2} \lambda_2^{\text{tree}} \|\boldsymbol{w}\|_2^2,$$

  where γ^{tree} penalizes tree complexity per leaf and $\lambda_1^{\text{tree}}, \lambda_2^{\text{tree}}$ apply ℓ_1 and ℓ_2 shrinkage to the leaf weights $\boldsymbol{w}$.

  – Every leaf is required to contain at least $n_{\min}$ observations ($n_j \geq n_{\min}$), which corresponds to the minimum leaf size in LightGBM and, under unit Hessian for squared-error loss, the minimum child weight in XGBoost.

  – **Random forest** (Breiman, 2001) forms the equally weighted average of B regression trees f_b, each grown independently on a bootstrap sample of the estimation window with

a random subset of candidate predictors evaluated at each split:

$$\hat{y}_{t,i} = \frac{1}{B}\sum_{b=1}^{B} f_b(\boldsymbol{x}_{t,i}).$$

5.3 Forecast evaluation

- **Forecast and realization.** We evaluate the unadjusted back-transformed variance forecast $\hat{\sigma}_t^2$ defined in Section 5.1 against the actual, unwinsorized realized variance of the closing interval, $\mathrm{RV}_t \equiv \mathrm{RV}_{t,c}$ (winsorization is applied to $\tilde{y}_{t,c}$ during model estimation only).
- **Evaluation set.** Because every forecast is generated by a model estimated on a strictly trailing window, all predictions are out of sample. We let $\mathcal{T}_{\mathrm{oos}}$ denote the set of common forecast-origin trading days following the initial estimation window and hyperparameter presample on which all competing forecasts are available, excluding early-close days, and write $T_{\mathrm{oos}} = |\mathcal{T}_{\mathrm{oos}}|$ (the exact sample dates and T_{oos} are given in Section 6.1).
- **Primary loss criterion.** Our primary statistical loss is the daily quasi-likelihood (QLIKE) loss of Patton (2011), which is robust to measurement noise in the realized-variance proxy. Because the diurnal factor $s_{t,c}^2$ cancels between $\mathrm{RV}_t = y_t^2 s_{t,c}^2$ (with unwinsorized y_t) and $\hat{\sigma}_t^2 = \hat{y}_t^2 s_{t,c}^2$, the daily loss takes the form

$$L_t = \frac{y_t^2}{\hat{y}_t^2} - \log\frac{y_t^2}{\hat{y}_t^2} - 1 = \mathrm{QLIKE}(\mathrm{RV}_t,\, \hat{\sigma}_t^2),$$

  with sample average over $\mathcal{T}_{\mathrm{oos}}$ given by

$$\overline{\mathrm{QLIKE}} = \bar{L} = \frac{1}{T_{\mathrm{oos}}}\sum_{t\in\mathcal{T}_{\mathrm{oos}}} L_t.$$

- **Benchmark forecast.** Our baseline is the closing-interval HAR benchmark: a HAR-type model in the spirit of Corsi (2009), estimated by OLS on the trailing averages $\bar{y}^{(k)}$ of the target over $k \in \{1, 5, 25, 125, 625, 3125\}$ samples and the 16 calendar columns $\boldsymbol{D}$ (Section 5.2), using closing-interval observations only (Appendix G.3).
- **Relative loss.** When comparing a candidate forecast A against a reference forecast B (the closing-interval HAR benchmark unless noted otherwise), we report the percentage difference in mean QLIKE:

$$\Delta\% = 100\Big(\frac{\bar{L}^A}{\bar{L}^B} - 1\Big),$$

  and assess its statistical significance via the Diebold–Mariano test detailed in Appendix I.
- **Squared-error loss.** In variance units, we also compute the mean squared error over $\mathcal{T}_{\mathrm{oos}}$:

$$\mathrm{MSE} = \frac{1}{T_{\mathrm{oos}}}\sum_{t\in\mathcal{T}_{\mathrm{oos}}} \big(\mathrm{RV}_t - \hat{\sigma}_t^2\big)^2.$$

- **Out-of-sample R^2.** Following Campbell and Thompson (2008), we compute R^2_{OOS} relative to the expanding historical mean of closing-interval realized variance known at the forecast origin:

$$R^2_{\mathrm{OOS}} = 1 - \frac{\sum_{t\in\mathcal{T}_{\mathrm{oos}}}\big(\mathrm{RV}_t - \hat{\sigma}_t^2\big)^2}{\sum_{t\in\mathcal{T}_{\mathrm{oos}}}\big(\mathrm{RV}_t - \overline{\mathrm{RV}}_{t-1}\big)^2}, \qquad \overline{\mathrm{RV}}_{t-1} = \frac{1}{|\{t' < t\}|}\sum_{t'<t}\mathrm{RV}_{t'},$$

where the expanding average $\overline{\mathrm{RV}}_{t-1}$ is taken over all prior trading days from the sample start specified in Section 6.1.

E33. Decision: to decide: compute MAE and the calibration diagnostics (the ratio $\overline{\mathrm{RV}}/\overline{\hat{\sigma}^2}$ and the under-prediction frequency) for the forecasts of Table 1, in make_rehaul_oos_rerun.py, or drop them.
Problem: MAE and calibration diagnostics promised here are not yet computed/reported.
No generator of the paper (writeup/make_rehaul_*.py) computes MAE, the ratio or the under-prediction frequency. The calibration files under results/atm_straddle_0dte_1530/ (recal_grid/, recal_2026-09-24/) come from the earlier recalibration studies, not from the current forecasts, and generated/campaign_numbers.tex holds one MAE of an earlier campaign, which this paper does not load.

6 Experiments and results

RR. Decision: a rerun of the pooled intraday models is approved (22 October 2026) and waits for the window choice.
Problem: the pooled rows of Table 1 and every pooled number in this section come from forecasts with a 1,000-day window that predate the late-start, no-observation and KKT-check fixes and still include the FOMC statement dummy; they will move. See the notes to the authors, "Rerun pending".

6.1 Experimental setup

- **Estimation samples** (see Appendix Table 4). The pooled intraday sample comprises 300,317 half-hour intervals beginning in June 1998,

$$300{,}317 = 303{,}442 \text{ intervals on the grid} - 3{,}125 \text{ burn-in samples};$$

  the pooled linear models and tree ensembles of Table 1 and of this section and Section 7 are trained on its 13 regular-hours bars with rolling windows of $W = 13{,}000$ bars (1,000 trading days; training scheme below), and the pooled linear runs of the model set on all its intervals with $W = 96{,}000$ intervals (Appendix B.2).
- **Evaluation span $\mathcal{T}_{\text{oos}}$** (see Section 5.3). We evaluate the rolling out-of-sample forecasts over the period from 3 January 2020 to 30 April 2024, which comprises $T_{\text{oos}} = |\mathcal{T}_{\text{oos}}| = 1{,}082$ trading days on which all competing forecasts are available (after excluding 6 early-close sessions). The pooled linear models and tree ensembles of Table 1 are scored on a longer span, from 3 January 2012 to 30 April 2024 (3,075 trading days, after excluding 26 early-close sessions), with a rolling window of 1,000 trading days; the linear penalties are chosen once on 4 validation folds covering 2008–2011, before the first forecast, and then frozen, and each tree ensemble keeps one configuration from an Optuna search on the pre-sample.
- **Training scheme.** Under the pooled intraday estimation scheme, the linear specifications and the tree ensembles are fit across the 13 regular-hours bars of each trading day using a rolling window of $W = 13{,}000$ bars (1,000 trading days), re-estimated at every bar; the pooled linear runs of the model set use all $M = 48$ half-hour intervals with $W = 96{,}000$ intervals. The models of Table 1, the linear models and the tree ensembles, are fit on the 13 regular-hours bars with a rolling window of 13,000 bars (1,000 trading days).

E15. Decision: one hyperparameter rule for every class, frozen after a single pre-sample pick.
Problem (2026-10-22): for the pooled models the rule now holds; the note can be closed once confirmed.
Every pooled model of Table 1 and of Sections 6–7 is fit on the 13 regular-hours bars with a window of 1,000 trading days ($W = 13{,}000$ bars), re-estimated at every bar: the linear penalties are chosen once on 4 pre-sample validation folds covering 2008–2011 and then frozen, and each tree method keeps one Optuna configuration (notes E34, E36). The tree configurations were selected on $W = 24{,}000$ intervals and are used as selected (note E24). The pooled linear runs of the model set on all 48 intervals ($W = 96{,}000$) are no longer the pooled rows of these sections.

- **Benchmark forecast.** Our baseline benchmark is the closing-interval HAR benchmark, a HAR-type model in the spirit of Corsi (2009) on the target's trailing averages over $k \in \{1, 5, 25, 125, 625, 3125\}$ samples and the 16 calendar columns, estimated by OLS on closing-interval observations (Appendix G.3); Appendix G.5 examines all model classes trained analogously on the closing 30-minute interval alone.
- **Rolling re-estimation with frozen hyperparameters.** The pooled intraday models are re-estimated at every 30-minute interval (see Appendix B.2). Their hyperparameters are chosen once, on the fixed pre-sample, and then held fixed for the whole evaluation span, with no re-tuning (Appendix D).
- **Evaluation metric and inference.** We evaluate predictive accuracy via the QLIKE loss against actual realized variance (see Section 5.3) and conduct pairwise inference using Diebold–Mariano tests with the Harvey–Leybourne–Newbold small-sample adjustment and a Newey–West long-run variance estimator ($q_{\mathrm{NW}} = 6$ on $\mathcal{T}_{\mathrm{oos}}$).[1]

We refer the reader to Appendix B for comprehensive implementation details.

6.2 Out-of-sample accuracy

- **Overview of forecasting performance.** Table 1 reports the out-of-sample MSE, average QLIKE loss, and out-of-sample R^2_{OOS} for each pooled intraday model class under both the HAR benchmark and the full predictor set across the 3,075 trading days from 3 January 2012 to 30 April 2024. The corresponding metrics for specifications estimated exclusively on the closing half-hour interval—including the OLS regression on the HAR benchmark predictor set, our closing-interval HAR benchmark forecast—are presented in Table 3 (see Appendix G.5). For the linear models, split at the start of 2020, the full predictor set has the lower QLIKE in both parts (2012–2019: 0.1718–0.1740 against 0.1840 for the HAR benchmark; 2020–2024: 0.1869–0.1874 against 0.1951–0.1952), and a higher R^2_{OOS} on 2012–2019 (68.8–69.3% against 66.4%) but a lower one on 2020–2024 (38.1–45.7% against 66.4%), so its lower R^2_{OOS} over the whole span comes from 2020–2024.

[1] Appendix I provides the complete specification of these statistical tests.

E36. Decision (user, 2026-10-20): Table 1 is like for like. It holds the pooled linear rerun, and since 2026-10-22 the pooled tree ensembles fitted on its design, trained and scored on the 13 regular-hours bars (the HAR horizons of Section 4, FOMC series in, a rolling window of 1,000 trading days, penalties chosen once on 4 validation folds covering 2008–2011 and then frozen), with every bar scored from 3 January 2012 to 30 April 2024 (3,075 trading days, one panel). Its tests are Diebold–Mariano on the daily average QLIKE. The reruns with the longer window and the rerun on every interval of the day are not in the table; their numbers are in results/oos_rerun_2026-10-20/ and results/oos_span_check_2026-10-20/. The abstract, the introduction, Data, this section, the conclusion and the appendix now give both spans: Table 1 from 3 January 2012, the closing-interval models, Table 3 and the straddle sample from 3 January 2020. Tree ensembles were held off until they were refit on the cluster. They are now refit on this design (2026-10-22, note E34: results/pooled_trees_rerun13_2026-10-20/w1000, every bar from 3 January 2012), and the table holds them: LightGBM, XGBoost and the random forest on both predictor sets, one configuration per method from an Optuna search on the pre-sample, frozen, scored on the same 39,975 bars as the linear rows (asserted) and given the same tests. The prose numbers of Section 6.3 read their 16:00 forecasts over $\mathcal{T}_{\text{oos}}$.
Problem: the abstract, the introduction and the conclusion still carry the earlier table's framing: the 15:30 forecast of the closing half hour, six model classes with trees, and the `\rhrAfPool...` counts, which writeup/make_rehaul_results.py now computes from the 16:00 rows of the forecasts of this table (2026-10-22, note E34). The Table 4 note gives these models' window, $W = 13{,}000$ bars.
New tests, QLIKE, all features against the HAR benchmark of the same model, p-values in parentheses: ridge -7.86 (<0.001), lasso -8.07 (<0.001), elastic net -10.45 (<0.001), LightGBM -11.67 (<0.001), XGBoost -10.78 (<0.001), random forest -11.86 (<0.001). The like-for-like values of the prose numbers are being computed (branch insurance-selling-2026-09-19-prose-numbers).

- **Computation of R^2_{OOS}.** Following Campbell and Thompson (2008), the reference forecast in R^2_{OOS} is an expanding historical mean of realized variance computed from 5 January 1998 through day $t-1$. For the pooled models of Table 1, linear and tree ensembles, it is the mean realized variance of the same half-hour bar over every earlier trading day; for the closing-interval models of Table 3 it is the mean realized variance of the closing half hour. Both rely solely on information available prior to the forecast origin.

6.3 The full predictor set versus the HAR benchmark

- **Paired comparison within each model class.** To isolate the predictive contribution of the exogenous features within each intraday model class m ($\mathcal{I} = \{1, \ldots, M\}$), we form the daily loss differential $d_t = L_t^{\text{Full}} - L_t^{\text{HAR}}$, where $L_t^{\mathcal{X}}$ denotes the QLIKE loss of the forecast $f_{m,\mathcal{X},\{1,\ldots,M\}}$. We report the percentage change in average loss $\Delta\%$ defined in Section 5.3, along with the Diebold–Mariano statistic DM computed on d_t over $\mathcal{T}_{\text{oos}}$ using lag truncation $q_{\text{NW}} = 6$.
- **QLIKE: full predictor set against the HAR benchmark.** Augmenting the HAR benchmark with the full predictor set lowers the average out-of-sample QLIKE loss in 4 of the 6 intraday model classes, with 0 of these reductions statistically significant at the 5% level (forecasts of the closing half hour over $\mathcal{T}_{\text{oos}}$). Scored on every bar of Table 1, from 3 January 2012 to 30 April 2024, the full predictor set has the lower QLIKE in 6 of the 6 classes, 6 of them significantly at the 5% level. On the closing half hour:
  - Regarding the tree ensembles, the reductions are $\Delta\% = -1.6$ for LightGBM (DM $= -0.60$, $p = 0.55$), -1.7 for XGBoost (DM $= -0.62$, $p = 0.53$), and -4.2 for the random forest (DM $= -1.73$, $p = 0.083$); the reduction is statistically significant at the 5% level in 0 of the three classes.
  - Regarding the penalized linear models, the QLIKE changes are $\Delta\% = -0.9$ for the elastic net (DM $= -0.31$, $p = 0.76$), 2.6 for ridge regression (DM $= 0.61$, $p = 0.54$), and 2.2 for the lasso (DM $= 0.58$, $p = 0.56$); the change is a statistically significant reduction

at the 5% level in 0 of the three classes.

6.4 Model classes

- **Exogenous features and the predictability of realized variance, by model class.** Comparing across model classes, the average QLIKE losses of the pooled penalized linear models and of the pooled tree ensembles on each predictor set are:
  - Under the full predictor set, the linear models attain average QLIKE losses between 0.1194 and 0.1236 (0.1194 for the elastic net, 0.1236 for ridge, and 0.1229 for the lasso), and the tree ensembles 0.1114–0.1140.
  - Under the HAR benchmark, the linear specifications achieve average QLIKE losses of 0.1203–0.1205, and the tree ensembles 0.1133–0.1190.
- **Comparison against the closing-interval HAR benchmark.** Every pooled intraday specification incurs a higher average QLIKE loss than the closing-interval HAR benchmark ($\overline{\text{QLIKE}} = 0.1033$):
  - Under the full predictor set, the average QLIKE loss of the pooled linear models exceeds that of the closing-interval HAR benchmark by +15.6 to +19.7%, and the difference is statistically significant at the 5% level for 3 of the three; for the pooled tree ensembles, the excess loss ranges from +7.9 to +10.4%, and it is significant at the 5% level for 1 of the three.
  - Under the HAR benchmark, the excess QLIKE loss of the pooled intraday models ranges from +9.7 to +16.7%; the difference is statistically significant at the 5% level for 6 of the 6 models and at the 1% level for 4.
  - Incorporating the exogenous covariates narrows the accuracy gap relative to the closing-interval HAR benchmark in 4 of the 6 intraday model classes.
  - Because the closing-interval HAR benchmark is estimated solely on the closing 30-minute interval of each day, we extend this closing-interval training scheme to all model classes in Appendix G.5.

6.5 Variable importance

RR. Decision: the pooled linear variable importance is computed on the current design (151 monthly re-estimations, results/rehaul_fi_pooled13_2026-10-22: dense across variables, sparse within each) and goes in after the rerun of the pooled models; the tree analysis is to be decided.
Problem: this subsection is still a placeholder. See the notes to the authors, "Rerun pending".

Variable importance is not yet computed for the pooled intraday model; the closing-interval analysis is parked (note E19).

E19. Decision (user, 2026-10-17): compute feature importance for the pooled intraday model with a monthly refit cadence, the default; the user has not chosen between weekly and monthly.
Problem: the pooled intraday model has no feature-importance analysis, and the parked closing-interval analysis refit every 10 trading days (writeup/make_rehaul_fi.py:446).
User decision 2026-10-17: switch to a monthly feature-importance refit; requires a run of feature importance for the pooled intraday model. No monthly cadence has been run, so this section reports none.

Table 1: Forecast accuracy of the intraday models, 3 January 2012 to 30 April 2024 (1,000-day window, hyperparameters frozen)

	HAR benchmark			All features			
Model	MSE (10^{-10})	QLIKE	R^2_{OOS} (%)	MSE (10^{-10})	QLIKE	R^2_{OOS} (%)	DM
OLS	0.94	0.1879	66.4	–	–	–	–
Ridge	0.94	0.1879	66.4	1.52	0.1774	45.4	−7.86 (<0.001)
Lasso	0.94	0.1879	66.4	1.58	0.1773	43.4	−8.07 (<0.001)
Elastic net	0.94	0.1879	66.4	1.40	0.1787	49.7	−10.45 (<0.001)
Tree ensembles							
LightGBM	1.03	0.1881	62.9	1.23	0.1698	55.7	−11.67 (<0.001)
XGBoost	1.02	0.1876	63.4	1.11	0.1653	60.2	−10.78 (<0.001)
Random forest	1.01	0.1893	63.5	1.16	0.1765	58.5	−11.86 (<0.001)

Notes: Pooled intraday linear models and tree ensembles, one design for every row: the HAR components at horizons $k \in \{1, 5, 25, 125, 625, 3125\}$, the full predictor set with the FOMC series, a rolling window of 1,000 trading days, re-estimated at every bar (Section 5), trained and scored on the 13 regular-hours half-hour bars ending 10:00 to 16:00 ET. Hyperparameters are chosen once and then frozen: the linear penalties on 4 validation folds covering 2008–2011, before the first forecast; the tree ensembles' configurations, one per method, by an Optuna search on the pre-sample (Appendix B.3). Every bar is scored on the $T_{\text{oos}} = 3{,}075$ trading days from 3 January 2012 to 30 April 2024 (early closes dropped), 39,975 bars (bars with a non-positive realized variance dropped); each forecast $\hat{\sigma}^2_{t,i} = \hat{y}^2_{t,i}\, s^2_{t,i}$ is scored against the actual realized variance $\text{RV}_{t,i}$ of its interval. MSE (in units of 10^{-10}) and QLIKE are averages over the scored intervals. $R^2_{\text{OOS}} = 1 - \sum(\text{RV}_{t,i} - \hat{\sigma}^2_{t,i})^2 / \sum(\text{RV}_{t,i} - \overline{\text{RV}}_{t-1,i})^2$, in percent, where $\overline{\text{RV}}_{t-1,i}$ is the historical mean realized variance of the same half-hour interval over every earlier trading day from 5 January 1998 (Campbell and Thompson, 2008). DM: the Diebold–Mariano statistic, with the correction of Harvey et al. (1997), for the daily average QLIKE of the full predictor set minus that of the HAR benchmark of the same model (negative: the full set has the lower loss), Newey–West lag 8, p-value in parentheses. Models trained on the last 30 minutes are in Table 3.

7 Economic value of the forecasts

RR. Decision: a rerun of the pooled intraday models is approved (22 October 2026) and waits for the window choice.
Problem: every pooled number in this section (the 12 pooled forecasts, their Sharpe ratios and tests) comes from forecasts that predate the late-start, no-observation and KKT-check fixes and the window choice, and will move. The closing-interval models of Section G.5 are not rerun. See the notes to the authors, "Rerun pending".

7.1 The straddle strategy

- **Option-trading sample.** We evaluate the strategy on the option-trading subsample $\mathcal{T}_{\text{opt}} \subset \mathcal{T}_{\text{oos}}$ (Section 6.1), the trading days with a same-day SPXW expiration:

$$T_{\text{opt}} = 866 \text{ trading days} = 872 \text{ SPXW expirations} - 6 \text{ early-close days}.$$

  The 419 subsequent SPXW expiration dates (413 full trading days) are held out from the present analysis.
- **Position.** We trade same-day-expiring (0DTE) SPXW options. With S_t denoting the underlying print on the option chain at the 15:30 ET forecast origin, we select the call and put at the nearest adjacent strikes with live quotes:

$$K_C = \min\{K \geq S_t\}, \qquad K_P = \max\{K \leq S_t\}.$$

  When $K_C = K_P = S_t$, the package forms an at-the-money straddle; otherwise, it forms a narrow strangle on the two strikes bracketing S_t. We refer to both positions as the straddle: an approximately delta-neutral package entered at 15:30 ET and held unhedged until expiration at 16:00 ET.

  E18. Decision: write the day filter into the text, and give the implied-volatility line the recovery used on the censored days.
  Problem: the day-drop rule (the wider strike more than 10 index points from S_t, or fewer than 10 quoted contracts) is still only a comment under this item, and the implied-volatility line and Appendix C make a censored leg missing, while the scored days recover it.
  S_t is now the 15:30 underlying print on the option chain and S_t^{close} the official S&P 500 close. The notebook checks that the Yahoo closes match the settlement column on every scored day. The chain field is copied as printed, and the code never labels it cash or futures, so the text calls it neither. A day is dropped when the wider strike is more than 10 index points from S_t, or when fewer than 10 contracts are quoted. That rule is in Appendix C and is still only a comment under this item. Early-close days are a separate drop, already in the T_{opt} line of the Data section. The vendor field is an hourly standard deviation from a bisection on [0.0005, 0.025]. Dividing by $\sqrt{2}$ is the half-hour conversion, and that is what was scored. On the 5 scored days where a leg is a failed-solve node, the hourly volatility is recovered by inverting the Black–76 price of the straddle on the 15:30 underlying, then divided by $\sqrt{2}$. Appendix C says a censored leg makes the implied variance missing. That is the drop path, not the scored path.
- **Entry prices.** From the legs' bid quotes b and ask quotes a at 15:30 ET, we define the package mid-quote P_t, bid P_t^b, and ask P_t^a as

$$P_t = \tfrac{1}{2}(b_t^C + a_t^C) + \tfrac{1}{2}(b_t^P + a_t^P), \qquad P_t^b = b_t^C + b_t^P, \qquad P_t^a = a_t^C + a_t^P.$$

- **Settlement.** The package settles in cash at the official S&P 500 close S_t^{close} with no intraday delta hedging; the simple return on the mid-quote premium paid at entry is

$$X_t = (S_t^{\text{close}} - K_C)^+ + (K_P - S_t^{\text{close}})^+, \qquad R_t = \frac{X_t}{P_t} - 1 \;\geq\; -1.$$

- **Implied variance.** Let IV_t^C and IV_t^P denote the call and put implied volatilities at 15:30 ET, quoted on a per-hour basis. The implied volatility over the closing half hour is

$$\mathrm{IV}_t = \frac{\mathrm{IV}_t^C + \mathrm{IV}_t^P}{2\sqrt{2}},$$

  and IV_t^2 is the corresponding implied variance to expiration.
- **Forecast variance risk premium.** Following Carr and Wu (2009) and Bollerslev et al. (2009), the variance risk premium over the closing half hour is $\mathrm{IV}_t^2 - \mathbb{E}_t[\mathrm{RV}_t]$, which we forecast by $\mathrm{IV}_t^2 - \hat{\sigma}_t^2$, where $\hat{\sigma}_t^2$ is the model variance forecast from Section 5.1. Every trading rule below conditions on the forecast exclusively through the variance gap $\hat{\sigma}_t^2 - \mathrm{IV}_t^2$, taking a long straddle position when forecast variance exceeds implied variance.
- **Trading strategies.** A strategy chooses a daily straddle position q_t ($q_t > 0$ long, $q_t < 0$ short) and earns the daily return $r_t = q_t R_t$. Because the variance forecast enters solely through q_t, comparing performance across rules isolates the economic value of the forecast (Fleming et al., 2001). We consider four rules:
  - *Unconditional short-straddle benchmark:* a static short-variance position setting $q_t \equiv -1$;
  - *Sign strategy:* $q_t = \mathrm{sign}(\hat{\sigma}_t^2 - \mathrm{IV}_t^2)$, with ties mapped to $q_t = -1$;
  - *Robust clip:* $q_t = \mathrm{clip}((\hat{\sigma}_t^2 - \mathrm{IV}_t^2)/\mathrm{MAD}_t, -1, 1)$, where $\mathrm{MAD}_t = \mathrm{median}_{t'<t}|\hat{\sigma}_{t'}^2 - \mathrm{IV}_{t'}^2|$ is the expanding median absolute variance gap over earlier days $t' < t$ in $\mathcal{T}_{\mathrm{opt}}$;
  - *Percentile rank:* the sign of the gap scaled by the empirical percentile rank of its magnitude across the $t-1$ preceding days $t' < t$ in $\mathcal{T}_{\mathrm{opt}}$:

$$q_t = \mathrm{sign}(\hat{\sigma}_t^2 - \mathrm{IV}_t^2) \cdot \frac{1}{t-1} \sum_{t'<t} \mathbf{1}\{|\hat{\sigma}_{t'}^2 - \mathrm{IV}_{t'}^2| \le |\hat{\sigma}_t^2 - \mathrm{IV}_t^2|\}.$$

- **Sizing sample.** Every position q_t is constructed causally using days $t' < t$ only. We reserve the first 252 days of $\mathcal{T}_{\mathrm{opt}}$ to initialize the scaling factor MAD_t, and evaluate all trading rules—including the sign strategy—on the remaining 614 post-warmup trading days ($t > 252$).
- **Execution.** At mid-quotes, the strategy return is $r_t = q_t R_t$. At quoted bid and ask prices, buying the straddle pays the ask P_t^a and selling receives the bid P_t^b; because settlement occurs in cash at expiration, the bid–ask spread is incurred at entry only:

$$r_t^{\mathrm{ba}} = \begin{cases} |q_t|\,(X_t/P_t^a - 1), & q_t > 0, \\ |q_t|\,(1 - X_t/P_t^b), & q_t < 0. \end{cases}$$

- **Transaction costs.** Executing at the quoted bid or ask rather than the mid-quote imposes a median half-spread cost of 2.9% of the straddle premium. No additional commissions or exchange fees are deducted.

7.2 Performance measures

- **Annualized Sharpe ratio.** We compute the annualized Sharpe ratio from daily strategy returns, both at mid-quotes and at quoted bid/ask prices, over the 614 post-warmup days of $\mathcal{T}_{\mathrm{opt}}$ ($t > 252$):

$$\mathrm{SR} = \frac{1}{|\mathcal{Y}|} \sum_{y \in \mathcal{Y}} \mathrm{SR}_y, \qquad \mathrm{SR}_y = \sqrt{A_y}\, \frac{\bar{r}_y}{\mathrm{sd}_y(r)}, \qquad A_y = \frac{N_y}{c_y},$$

where $\mathcal{Y}$ is the set of the 4 calendar years of those 614 days, N_y is the number of them that fall in year y (option days, i.e. days with a same-day SPXW expiration), $\bar{r}_y$ and $\mathrm{sd}_y(r)$ are the sample mean and standard deviation (divisor $N_y - 1$) of r_t over those N_y days, and c_y is the fraction of calendar year y that the sample covers, from the later of 1 January and the sample's first day to the earlier of 31 December and its last day. A_y is thus the number of option days per calendar year at year y's expiration schedule: 163.4, 219.0, 248.0 and 251.1 for 2021 to 2024 (N_y = 64, 219, 248, 83; c_y = 0.392, 1.000, 1.000, 0.331). Each calendar year enters with equal weight and is annualized at its own trading frequency. The frequency changes within the sample: SPXW expirations were listed on Mondays, Wednesdays and Fridays until spring 2022 and on every weekday from 16 May 2022. Equal weights also keep a year with many option days from dominating the average if the strategy's performance changes over the sample. Table 2 (Appendix F) reports the main Sharpe ratio comparison under two other definitions: the per-year ratio with $\sqrt{N_y}$ in place of $\sqrt{A_y}$, the years weighted by N_y, and the ratio of the mean to the standard deviation of all 614 daily returns, scaled by $\sqrt{252}$.

D6. Decision (user, 2026-10-22): the Sharpe ratio is the definition above (D6): for each calendar year y of the post-warmup sample, $\mathrm{SR}_y = \sqrt{A_y}\,\bar{r}_y/\mathrm{sd}_y(r)$ with $A_y = N_y/c_y$ the year's option days per calendar year, and SR the equal-weight mean of the SR_y; every statement about a Sharpe ratio difference uses the paired block-bootstrap test of Appendix I, which resamples D6 itself.
Disclosure: this definition was adopted after the results under the previous one (the per-year ratio with $\sqrt{N_y}$, N_y-weighted, of 2026-10-18) had been seen. That ratio (D1) and the ratio of the pooled daily returns scaled by $\sqrt{252}$ (D2), with the Ledoit and Wolf (2008) test of the pooled ratios, are reported in Table 2, together with D6 before and from 16 May 2022.

N15b. Decision: to decide (rerun needed).
Problem: the interval-shift numbers (`shiftSwingMed*`, `shiftNoiseBand` = $1.96\sqrt{252/n}$; Section 7.3 and the appendix shift-test paragraph) come from experiments/close_shift_test_headline.py, which still annualizes with the fixed $\sqrt{252}$ on all option days. They were not regenerated: that script rewrites results/close_shift_test/ and its figures. The compounding growth rate G below keeps its factor 252: it annualizes a mean log growth rate, not a Sharpe ratio.

The 216 days of $\mathcal{T}_{\mathrm{oos}}$ without a same-day SPXW expiration enter the statistical accuracy evaluation only. Because the position is held for 30 minutes intraday, the risk-free return over the holding period is negligible and omitted.

- **Mean return and its t-statistic.** We test the statistical significance of mean daily returns over the 614 post-warmup days $\mathcal{T}^{+}_{\mathrm{opt}} = \{t \in \mathcal{T}_{\mathrm{opt}} : t > 252\}$, the sample of the Sharpe ratio, with the test that Lazarus et al. (2018) recommend: the t-statistic with the Newey–West (Bartlett kernel) long-run variance at bandwidth $S = \lceil 1.3\,(614)^{1/2} \rceil = 33$, whose autocovariances divide by the number of days with no small-sample factor, and fixed-b critical values (Kiefer and Vogelsang, 2005); the two-sided 5% critical value is 2.11 (Appendix I).
- **Hit rate and position counts.** We also record the empirical hit rate $\frac{1}{614}\sum_{t\in\mathcal{T}^{+}_{\mathrm{opt}}} \mathbf{1}\{r_t > 0\}$ and the number of long-straddle days $\#\{t \in \mathcal{T}^{+}_{\mathrm{opt}} : q_t > 0\}$.
- **Gain over the short-straddle benchmark.** For any strategy A compared against a benchmark strategy B (the unconditional short straddle unless noted otherwise), we define the daily return differential $\Delta_t = r^A_t - r^B_t$ and the Sharpe ratio gain

$$\Delta\mathrm{SR} = \mathrm{SR}(r^A) - \mathrm{SR}(r^B).$$

The t-statistic of the previous item, applied to Δ_t, tests whether the mean return differential differs from zero. Because it does not test the equality of the two Sharpe ratios, every statement about ΔSR uses a paired circular block bootstrap of ΔSR itself on the same days

(Appendix I); the studentized bootstrap test of Ledoit and Wolf (2008), which tests the equality of the Sharpe ratios of the pooled daily returns, is reported as a robustness check.

- **Elementary score.** To measure forecast accuracy specifically at the decision threshold IV_t^2, we evaluate the elementary score of Ehm et al. (2016) normalized by implied variance:

$$\mathrm{ES}_{1,t} = \left|\frac{\mathrm{RV}_t}{\mathrm{IV}_t^2} - 1\right| \mathbf{1}\{\min(\hat{\sigma}_t^2, \mathrm{RV}_t) \le \mathrm{IV}_t^2 < \max(\hat{\sigma}_t^2, \mathrm{RV}_t)\},$$

  and denote its sample mean over the post-warmup days $\mathcal{T}_{\mathrm{opt}}^+$ by $\overline{\mathrm{ES}}_1$.
- **Decision-threshold property.** Note that $\mathrm{ES}_{1,t}$ is strictly positive if and only if the forecast $\hat{\sigma}_t^2$ and the realization RV_t fall on opposite sides of the implied variance IV_t^2—that is, precisely when the sign strategy takes the wrong directional position.
- **Compounding.** When evaluating compounded wealth trajectories from initial capital $V_0 = 1$, we allocate a fixed fraction $f = 3\%$ of current wealth to the straddle premium each day, yielding terminal portfolio value and annualized log growth rate

$$V = \prod_{t \in \mathcal{T}_{\mathrm{opt}}^+} (1 + f\, r_t), \qquad G = \frac{252}{614} \sum_{t \in \mathcal{T}_{\mathrm{opt}}^+} \log(1 + f\, r_t).$$

- Under bid/ask execution, r_t^{ba} replaces $r_t = q_t R_t$ in the compounding recursion. We define maximum drawdown as $\max_t(1 - V_t / \max_{t' \le t} V_{t'})$, where V_t is the cumulative portfolio value at the end of day t.

7.3 Results

- **Unconditional short-straddle benchmark.** Selling the closing straddle unconditionally ($q_t \equiv -1$; labeled "Short straddle" in the tables) yields SR $= -0.12$ at mid-quotes and -0.46 at quoted bid/ask prices. It has a positive return on 62.2% of the post-warmup days, and its mean daily return is statistically indistinguishable from zero ($t = 0.67$, fixed-b $p = 0.52$), and the strategy loses money in 2 of the 4 calendar years: passively selling the closing 0DTE straddle earns no statistically detectable variance risk premium in our option-data subsample.
- **Sign strategies of the intraday models versus the short-straddle benchmark.** In point estimates, 12 of the 12 pooled intraday forecasts achieve a positive Sharpe ratio gain ($\Delta\mathrm{SR} > 0$) over the unconditional short straddle.

  N15c. Decision: to decide (James; options in writeup/RF_STORY_OPTIONS.md).
  Problem: under the Sharpe ratio of note D6, the full predictor set raises the pooled Sharpe ratio in 6 of the 6 classes (ΔSR from +0.23 to +0.97; elastic net +0.40, ridge +0.25, lasso +0.37, random forest +0.23), so "every intraday model class" (01_introduction, 09_conclusion, the abstract in paper_rehaul.tex) now holds in point estimates. The bootstrap test finds 0 of these gains statistically significant at the 5% level; those sentences state no significance. Under the previous ratio (D1) the count was 2 of the 6 classes (Table 2), and D6 was adopted after the D1 results had been seen (note D6). 12 of the 12 pooled forecasts have $\Delta\mathrm{SR} > 0$ over the short straddle (this item). The items of this section print counts and ranges.

  - The mean return gain over the short-straddle benchmark is statistically significant at the 5% level in 0 of the 12 specifications (fixed-b test of the mean daily return differential on the 614 post-warmup days, bandwidth $S = 33$); the t-statistics range from 0.79 to 1.70.
- **Sharpe ratios of the pooled intraday models.** On the HAR benchmark, the Sharpe ratios of the 6 pooled intraday specifications range from 0.59 to 1.07 at mid-quotes and from 0.23 to 0.71 at quoted bid/ask prices; on the full predictor set, they range from 1.14 to 1.56 at mid-quotes and from 0.79 to 1.21 at quoted bid/ask prices.

- **Full predictor set versus the HAR benchmark.** Expanding the information set from the HAR benchmark to the full predictor set increases the mid-quote Sharpe ratio in 6 of the 6 intraday model classes; across classes, the mid-quote Sharpe ratio ranges from 0.59 to 1.07 on the HAR benchmark and from 1.14 to 1.56 on the full predictor set.
  - The resulting gain ΔSR spans +0.23 to +0.97 across classes, with a median of +0.38 (elastic net +0.40, ridge +0.25, lasso +0.37).
  - The mean daily return differential is statistically significant at the 5% level in 0 of the 6 classes (fixed-b test, bandwidth $S = 33$).
  - The Sharpe ratio gain is statistically significant at the 5% level in 0 of the 6 classes in which it is positive (block-bootstrap test of equal Sharpe ratios, Appendix I); the Ledoit and Wolf (2008) test of the Sharpe ratios of the pooled daily returns rejects their equality at the 5% level in 0 of the 6 classes (Table 2).
- **Position sizing.** Across the 12 intraday forecasts, the median mid-quote Sharpe ratio is 1.11 under the sign strategy, 0.88 under the robust clip rule, and 0.95 under the percentile-rank rule.
  - In point estimates, continuous sizing via the robust clip outperforms the binary sign strategy ($\mathrm{SR}^{\mathrm{clip}} > \mathrm{SR}^{\mathrm{sign}}$) for 4 of the 12 intraday forecasts.
- **Concentration of economic gains.** As is typical of short-dated option strategies, cumulative performance is concentrated in a small subset of trading days:
  - removing each forecast's 10 best days renders the mean return differential $\bar{\Delta}$ over the short-straddle benchmark negative for 12 of the 12 intraday forecasts;
  - trading days on which the straddle expires worthless ($R_t = -1$, corresponding to settlement pinning near the strike; 137 of the 614 post-warmup days) account for 39% of the positive returns of the median pooled intraday forecast (36–40% across the 12 forecasts).
- **Statistical accuracy versus economic value.** Across the 12 pooled intraday forecasts evaluated on $\mathcal{T}_{\mathrm{opt}}$, the Spearman rank correlations between average statistical loss and the post-warmup Sharpe ratio (where a negative correlation indicates that lower loss corresponds to higher SR) are

$$\rho_S(\overline{\mathrm{QLIKE}}, \mathrm{SR}) = -0.13, \qquad \rho_S(\overline{\mathrm{ES}}_1, \mathrm{SR}) = -0.54.$$

  Among the pooled intraday forecasts, the rank correlation of the Sharpe ratio is -0.13 with average QLIKE loss and -0.54 with $\overline{\mathrm{ES}}_1$, the loss at the implied-variance threshold. With the per-year ratio scaled by $\sqrt{N_y}$ and weighted by N_y in place of SR, the two correlations are 0.05 and 0.31; with the ratio of the pooled daily returns scaled by $\sqrt{252}$, they are -0.12 and -0.06 (note D6).
- **Robustness checks** (detailed in Appendix F):
  - *Interval-shift test:* when we replace $\hat{y}_{t,c}$ in the sign rule by the forecast $\hat{y}_{t,c+j}$ from j half-hour intervals away ($j < 0$ earlier on day t; $j > 0$ on day $t+1$), $q_t = \mathrm{sign}(\hat{y}^2_{t,c+j}\, s^2_{t,c} - \mathrm{IV}^2_t)$, the median Sharpe ratio across the 9 earlier-draft forecasts that hold every intraday interval (note N18; 865 option-data days, annualized with a fixed factor, note N15b) varies between 0.00 and 1.58 across j (a half-range of ± 0.79). Because this variation lies well within the ± 1.06 95% sampling band of an annualized Sharpe ratio under i.i.d. returns, differences of this magnitude are consistent with sampling variation;

E26. Decision: to decide: restate on the pooled intraday forecasts only, or move to the closing-interval appendix.
Problem: the bullet reports a median Sharpe ratio across models without naming them, and the forecasts it shifts are not the pooled intraday forecasts of the tables.
The shift-test macros come from experiments/close_shift_test_headline.py through generated/close_shift_test_numbers.tex. Per results/close_shift_test/SUMMARY.md and shift_test_modelset.csv, the shifted forecasts are the "paper" family of the model set (a0, blk2, blk2_inc, lgbm, xgb, lasso_t, lasso_f, enet, b2lasso): the earlier draft's pooled 48-bar intraday forecasts. The current pooled linear rows (poolpt_*) and the per-bar rows are not shifted, because their tables hold only the 16:00 row (shift_test_not_shifted.csv). For $j > 0$ the shifted forecast is the one of day $t + 1$, which is not in $\mathcal{F}_t$ at the 15:30 forecast origin, so those offsets are a placebo, and the bullet does not say so.

– *Subsample stability:* we verify performance across subsamples split around May 2022 (when CBOE expanded SPXW expirations to all five weekdays; Table 2);
– *Compounded wealth:* across all 112 forecasts in the model set, compounding at $f = 3\%$ of wealth per day over the 614 post-warmup days yields a median terminal portfolio value V of 2.4 at mid-quotes and 1.4 at quoted bid/ask prices, with maximum drawdowns reaching up to 59% and 68%, respectively.

E27. Decision: restrict the aggregation to intraday forecasts only (the 12 pooled intraday forecasts of note E25), dropping the closing-interval forecasts.
Problem: compounded wealth is computed across all 112 forecasts of the model set, closing-interval forecasts included, and the generator cannot yet compute it on the intraday forecasts of note E25.
The macros come from experiments/close_compounding_headline.py, which reads per-day returns from the master table (results/close_master_table/master_table_daily.parquet) for the forecasts of paper_model_set.csv. That set holds the per-bar trees and per-bar linear forecasts of the closing-interval scheme, the six pooled linear rows (poolpt_*), and the earlier draft's pooled 48-bar forecasts (the "paper" family, whose lgbm and xgb rows are the earlier draft's trees). None of the six pooled trees is in it: they are the base-5 refit of note E34 (results/pooled_trees_rerun13_2026-10-20/w1000, read at the 16:00 bar by writeup/pooled_trees_13.py), so a group or prefix filter cannot select the 12 forecasts. Missing: their per-day straddle returns in the compounding input, or a loader for that table in the compounding script (make_rehaul_economics.py already builds those returns). The pooled-linear group alone is already in results/close_compounding/SUMMARY.md, with no macro.

Appendix G.5 reports each model class trained on the closing half hour alone.

8 Conclusion

- **Exogenous predictors and QLIKE loss in pooled intraday models.** Relative to the HAR benchmark, augmenting the model with the full exogenous predictor set lowers mean QLIKE in 4 of the 6 pooled intraday model classes, with 0 of these reductions statistically significant at the 5% level (Section 6). The reduction is statistically significant at the 5% level in 0 of the three penalized linear classes and in 0 of the three tree classes.
- **Comparison with the closing-interval HAR benchmark.** Every pooled intraday specification incurs a higher mean QLIKE than the closing-interval HAR benchmark, the OLS HAR regression estimated solely on closing-interval data. Incorporating the full exogenous predictor set narrows this loss gap in 4 of the 6 model classes; on the full predictor set the gap is statistically significant at the 5% level for 3 of the three penalized linear models and for 1 of the three tree ensembles (Section 6.4).
- **Models trained on the closing half hour.** Results for each model class estimated on the closing interval alone are reported in Appendix G.5.
- **Economic value in 0DTE straddles.** When the variance forecasts are used to trade same-day-expiring SPXW straddles over the closing half hour on the option-data subsample, adding the exogenous predictors increases the sign strategy's post-warmup Sharpe ratio in every intraday model class. Across the 12 pooled intraday forecasts, the rank correlation of the straddle Sharpe ratio is -0.13 with average QLIKE loss and -0.54 with $\overline{\mathrm{ES}}_1$, the loss at the implied-variance threshold (Section 7.3).

  N15c. Problem: under the Sharpe ratio of note D6, the full predictor set raises the pooled Sharpe ratio in 6 of the 6 classes (ΔSR from $+0.23$ to $+0.97$), so "every intraday model class" holds in point estimates; 0 of these gains are statistically significant at the 5% level (bootstrap test), and under the previous ratio the count was 2 of 6. See notes N15c and D6 in Section 7.

- **Scope and limitations.**
  - Our out-of-sample evaluation spans $T_{\mathrm{oos}} = 1{,}082$ trading days from 3 January 2020 to 30 April 2024; the forecast accuracy of the pooled linear models and tree ensembles (Table 1) is scored on 3,075 trading days from 3 January 2012 to 30 April 2024.
  - The empirical analysis focuses on a single equity index (the S&P 500) and a single target horizon (the 15:30–16:00 ET closing half hour).
  - In our economic evaluation, the straddle is held unhedged to cash settlement, and executing at quoted bid and ask prices rather than mid-quotes materially reduces the strategy's Sharpe ratio (Section 7).
- **Directions for future work.** Natural extensions include out-of-sample evaluation on post-30 April 2024 option expirations, cross-market validation on additional index option chains, and intraday variance forecasting at earlier decision origins prior to 15:30 ET.

References

Takuya Akiba, Shotaro Sano, Toshihiko Yanase, Takeru Ohta, and Masanori Koyama. Optuna: A next-generation hyperparameter optimization framework. In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 2623–2631, 2019.

Caio Almeida, Gustavo Freire, and Rodrigo Hizmeri. 0DTE asset pricing. `https://ssrn.com/abstract=4701401`, 2025. Working paper, SSRN 4701401.

Diego Amaya, Peter Christoffersen, Kris Jacobs, and Aurelio Vasquez. Does realized skewness predict the cross-section of equity returns? *Journal of Financial Economics*, 118(1):135–167, 2015. doi: 10.1016/j.jfineco.2015.02.009.

Torben G. Andersen and Tim Bollerslev. Intraday periodicity and volatility persistence in financial markets. *Journal of Empirical Finance*, 4(2-3):115–158, 1997.

Torben G. Andersen, Tim Bollerslev, Francis X. Diebold, and Paul Labys. Modeling and forecasting realized volatility. *Econometrica*, 71(2):579–625, 2003.

Torben G. Andersen, Tim Bollerslev, Peter F. Christoffersen, and Francis X. Diebold. Volatility and correlation forecasting. In *Handbook of Economic Forecasting*, volume 1, pages 777–878. Elsevier, 2006.

Torben G. Andersen, Tim Bollerslev, and Francis X. Diebold. Roughing it up: including jump components in the measurement, modeling, and forecasting of return volatility. *The Review of Economics and Statistics*, 89(4):701–720, 2007.

Donald W. K. Andrews. Heteroskedasticity and autocorrelation consistent covariance matrix estimation. *Econometrica*, 59(3):817–858, 1991. doi: 10.2307/2938229.

Federico M. Bandi, Nicola Fusari, and Roberto Renò. 0DTE option pricing. `https://ssrn.com/abstract=4503344`, 2023. Working paper, SSRN 4503344.

Ole E. Barndorff-Nielsen and Neil Shephard. Econometric analysis of realized volatility and its use in estimating stochastic volatility models. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 64(2):253–280, 2002. doi: 10.1111/1467-9868.00336.

Ole E. Barndorff-Nielsen and Neil Shephard. Power and bipower variation with stochastic volatility and jumps. *Journal of Financial Econometrics*, 2(1):1–37, 2004.

Ole E. Barndorff-Nielsen, Silja Kinnebrock, and Neil Shephard. Measuring downside risk: Realized semivariance. In Tim Bollerslev, Jeffrey R. Russell, and Mark W. Watson, editors, *Volatility and Time Series Econometrics: Essays in Honor of Robert Engle*, pages 117–136. Oxford University Press, 2010. doi: 10.1093/acprof:oso/9780199549498.003.0007.

Heiner Beckmeyer, Nicole Branger, and Leander Gayda. Retail traders love 0DTE options... but should they? `https://ssrn.com/abstract=4404704`, 2023. Working paper, SSRN 4404704.

Jeremias Bekierman and Hans Manner. Forecasting realized variance measures using time-varying coefficient models. *International Journal of Forecasting*, 34(2):276–287, 2018.

Mikkel Bennedsen, Asger Lunde, and Mikko S. Pakkanen. Decoupling the short- and long-term behavior of stochastic volatility. *Journal of Financial Econometrics*, 20(5):961–1006, 2022.

Fischer Black and Myron Scholes. The pricing of options and corporate liabilities. *Journal of Political Economy*, 81(3):637–654, 1973.

Tim Bollerslev, George Tauchen, and Hao Zhou. Expected stock returns and variance risk premia. *Review of Financial Studies*, 22(11):4463–4492, 2009. doi: 10.1093/rfs/hhp008.

Tim Bollerslev, Andrew J. Patton, and Rogier Quaedvlieg. Exploiting the errors: a simple approach for improved volatility forecasting. *Journal of Econometrics*, 192(1):1–18, 2016.

Tim Bollerslev, Benjamin Hood, John Huss, and Lasse Heje Pedersen. Risk everywhere: modeling and managing volatility. *The Review of Financial Studies*, 31(7):2729–2773, 2018.

Leo Breiman. Random forests. *Machine Learning*, 45(1):5–32, 2001. doi: 10.1023/A:1010933404324.

Jonathan Brogaard, Jaehee Han, and Peter Y. Won. Does 0DTE options trading increase volatility? `https://ssrn.com/abstract=4426358`, 2023. Working paper, SSRN 4426358.

Andrea Bucci. Realized volatility forecasting with neural networks. *Journal of Financial Econometrics*, 18(3):502–531, 2020. doi: 10.1093/jjfinec/nbaa008.

John Y. Campbell and Samuel B. Thompson. Predicting excess stock returns out of sample: Can anything beat the historical average? *Review of Financial Studies*, 21(4):1509–1531, 2008. doi: 10.1093/rfs/hhm055.

Peter Carr and Liuren Wu. Variance risk premiums. *Review of Financial Studies*, 22(3):1311–1341, 2009. doi: 10.1093/rfs/hhn038.

Tianqi Chen and Carlos Guestrin. XGBoost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 785–794, 2016.

Kim Christensen, Mathias Siggaard, and Bezirgen Veliyev. A machine learning approach to volatility forecasting. *Journal of Financial Econometrics*, 21(5):1680–1727, 2023.

Fulvio Corsi. A simple approximate long-memory model of realized volatility. *Journal of Financial Econometrics*, 7(2):174–196, 2009.

Fulvio Corsi, Stefan Mittnik, Christian Pigorsch, and Uta Pigorsch. The volatility of realized volatility. *Econometric Reviews*, 27(1-3):46–78, 2008.

Francis X. Diebold and Roberto S. Mariano. Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 13(3):253–263, 1995.

Chukwuma Dim, Bjorn Eraker, and Grigory Vilkov. 0DTEs: Trading, gamma risk and volatility propagation. `https://ssrn.com/abstract=4692190`, 2024. Working paper, SSRN 4692190.

Werner Ehm, Tilmann Gneiting, Alexander Jordan, and Fabian Krüger. Of quantiles and expectiles: Consistent scoring functions, Choquet representations and forecast rankings. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 78(3):505–562, 2016.

Robert F. Engle and Giampiero M. Gallo. A multiple indicators model for volatility using intra-daily data. *Journal of Econometrics*, 131(1-2):3–27, 2006.

Robert F. Engle and Magdalena E. Sokalska. Forecasting intraday volatility in the US equity market. multiplicative component GARCH. *Journal of Financial Econometrics*, 10(1):54–83, 2012. doi: 10.1093/jjfinec/nbr005.

Jeff Fleming, Chris Kirby, and Barbara Ostdiek. The economic value of volatility timing. *Journal of Finance*, 56(1):329–352, 2001. doi: 10.1111/0022-1082.00327.

Jerome Friedman, Trevor Hastie, and Robert Tibshirani. Regularization paths for generalized linear models via coordinate descent. *Journal of Statistical Software*, 33(1):1–22, 2010. doi: 10.18637/jss.v033.i01.

Jerome H. Friedman. Greedy function approximation: A gradient boosting machine. *Annals of Statistics*, 29(5):1189–1232, 2001. doi: 10.1214/aos/1013203451.

Pierre Garrigues and Laurent El Ghaoui. An homotopy algorithm for the lasso with online observations. In *Advances in Neural Information Processing Systems*, volume 21, 2008.

Domenico Giannone, Michele Lenza, and Giorgio E. Primiceri. Economic predictions with big data: The illusion of sparsity. *Econometrica*, 89(5):2409–2437, 2021.

Shihao Gu, Bryan Kelly, and Dacheng Xiu. Empirical asset pricing via machine learning. *Review of Financial Studies*, 33(5):2223–2273, 2020.

Peter R. Hansen. A test for superior predictive ability. *Journal of Business & Economic Statistics*, 23(4):365–380, 2005. doi: 10.1198/073500105000000063.

Peter R. Hansen, Asger Lunde, and James M. Nason. The model confidence set. *Econometrica*, 79 (2):453–497, 2011.

Peter R. Hansen, Zhuo Huang, and Howard Howan Shek. Realized garch: a joint model for returns and realized measures of volatility. *Journal of Applied Econometrics*, 27(6):877–906, 2012.

David Harvey, Stephen Leybourne, and Paul Newbold. Testing the equality of prediction mean squared errors. *International Journal of Forecasting*, 13(2):281–291, 1997.

Rodrigo Hizmeri and Mattia Bevilacqua. Early birds get the vol: Morning volatility uncertainty and variance risk premium. `https://ssrn.com/abstract=6212458`, 2025. Working paper, SSRN 6212458.

Arthur E. Hoerl and Robert W. Kennard. Ridge regression: Biased estimation for nonorthogonal problems. *Technometrics*, 12(1):55–67, 1970.

Guolin Ke, Qi Meng, Thomas Finley, Taifeng Wang, Wei Chen, Weidong Ma, Qiwei Ye, and Tie-Yan Liu. LightGBM: A highly efficient gradient boosting decision tree. In *Advances in Neural Information Processing Systems*, volume 30, 2017.

Nicholas M. Kiefer and Timothy J. Vogelsang. A new asymptotic theory for heteroskedasticity-autocorrelation robust tests. *Econometric Theory*, 21(6):1130–1164, 2005. doi: 10.1017/S0266466605050565.

Serhiy Kozak, Stefan Nagel, and Shrihari Santosh. Shrinking the cross-section. *Journal of Financial Economics*, 135(2):271–292, 2020.

Eben Lazarus, Daniel J. Lewis, James H. Stock, and Mark W. Watson. HAR inference: Recommendations for practice. *Journal of Business & Economic Statistics*, 36(4):541–559, 2018. doi: 10.1080/07350015.2018.1506926.

Oliver Ledoit and Michael Wolf. Robust performance hypothesis testing with the Sharpe ratio. *Journal of Empirical Finance*, 15(5):850–859, 2008. doi: 10.1016/j.jempfin.2008.03.002.

Robert C. Merton. Theory of rational option pricing. *Bell Journal of Economics and Management Science*, 4(1):141–183, 1973.

Whitney K. Newey and Kenneth D. West. A simple, positive semi-definite, heteroskedasticity and autocorrelation consistent covariance matrix. *Econometrica*, 55(3):703–708, 1987. doi: 10.2307/1913610.

Whitney K. Newey and Kenneth D. West. Automatic lag selection in covariance matrix estimation. *Review of Economic Studies*, 61(4):631–653, 1994. doi: 10.2307/2297912.

Andrew J. Patton. Volatility forecast comparison using imperfect volatility proxies. *Journal of Econometrics*, 160(1):246–256, 2011.

Eghbal Rahimikia and Ser-Huang Poon. Machine learning for realised volatility forecasting. `https://ssrn.com/abstract=3707796`, 2020. Working paper, SSRN 3707796.

Joseph P. Romano and Michael Wolf. Stepwise multiple testing as formalized data snooping. *Econometrica*, 73(4):1237–1282, 2005. doi: 10.1111/j.1468-0262.2005.00615.x.

Zhouyu Shen and Dacheng Xiu. Can machines learn weak signals? Technical report, NBER Working Paper 33421, 2025.

Stephen J. Taylor and Xinzhong Xu. The incremental volatility information in one million foreign exchange quotations. *Journal of Empirical Finance*, 4(4):317–340, 1997. doi: 10.1016/S0927-5398(97)00010-8.

Robert Tibshirani. Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society: Series B*, 58(1):267–288, 1996.

Mark A. van de Wiel, Mirrelijn M. van Nee, and Armin Rauschenberger. Fast cross-validation for multi-penalty high-dimensional ridge regression. *Journal of Computational and Graphical Statistics*, 30(4):835–847, 2021. doi: 10.1080/10618600.2021.1904962.

Wessel N. van Wieringen. Lecture notes on ridge regression, 2023. Version 0.60 (v8), 27 June 2023.

Grigory Vilkov. 0DTE trading rules: Tail risk, implementation, and tactical timing. `https://ssrn.com/abstract=4641356`, 2026. Working paper, SSRN 4641356; first posted 2023 as "0DTE Trading Rules".

Chao Zhang, Yihuang Zhang, Mihai Cucuringu, and Zhongmin Qian. Volatility forecasting with machine learning and intraday commonality. *Journal of Financial Econometrics*, 22(2):492–530, 2024. doi: 10.1093/jjfinec/nbad005.

Hui Zou. The adaptive lasso and its oracle properties. *Journal of the American Statistical Association*, 101(476):1418–1429, 2006. doi: 10.1198/016214506000000735.

Hui Zou and Trevor Hastie. Regularization and variable selection via the elastic net. *Journal of the Royal Statistical Society: Series B*, 67(2):301–320, 2005.

A Appendix structure

- **Implementation details** (Appendix B): samples and windows, the training schemes in practice, the tuned hyperparameters and the tuning procedures.
- **Data and predictors** (Appendix C):

  - sources and spans;
  - predictor dictionary: 46 exogenous series and 586 intraday predictors in the full predictor set (47 series and 628 predictors in the closing-interval model as fitted, which still carries the FOMC statement-interval dummy dropped from the dictionary; note E35) (Appendix Table 6);
  - missing and stale data;
  - removal of duplicate predictors; the window filter (constant and duplicate predictors dropped in each estimation window $\mathcal{W}_t$);
  - SPXW option chain; implied-volatility inversion;
  - FOMC calendar.
- **Hyperparameter selection by model** (Appendix D):
  - linear models: penalty strength λ, elastic-net mixing α; block-penalized ridge, penalty weights ω_j;
  - tree ensembles: fixed hyperparameters, Optuna budgets, selected trials; boosting rounds $b = 1, \ldots, B$, learning rate η;
  - re-selection frequency, on the estimation sample $\mathcal{W}_t^{\text{est}}$, gap $\mathcal{W}_t^{\text{gap}}$ and validation sample $\mathcal{W}_t^{\text{val}}$; re-estimation frequency;
- **Robustness, statistical** (Appendix E):
  - full model table, tree ensembles included (not reported);
  - Diebold–Mariano and HAC conventions (DM with the Harvey–Leybourne–Newbold correction and Newey–West lag truncation q_{NW}; fixed-b tests of mean returns; the block-bootstrap test of equal Sharpe ratios and the Ledoit–Wolf test of the pooled ratios);
  - model confidence set, superior predictive ability, StepM;
  - regime and break tests (not reported);
  - closing-interval vs intraday models;
  - remaining variance through the trading day (not reported).
- **Robustness, trading** (Appendix F):
  - date-shift test;
  - Murphy diagrams (elementary scores); subsample splits;
  - P&L decomposition (not reported);
  - compounding;
  - position sizing;
  - strategy variations (not reported); other forecast origins (not reported); month-ends (not reported).
- **Closing-interval models** (Appendix G): the closing-interval sample (Appendix G.1), predictors and column counts (Appendix G.2), the closing-interval HAR benchmark and its estimation (Appendix G.3), closing-interval standardization (Appendix G.4), and every model class trained on the last 30 minutes (Appendix G.5).
- **Reproducibility** (Appendix H):
  - software lock files (not reported);
  - hashes of the design matrices (not reported);
  - computing environment;
  - CPU classes (not reported);
  - compute time (not reported);

– the generator script behind every reported number.

- **Statistical tests** (Appendix I): the Diebold–Mariano statistic, the Newey–West variance, the fixed-b test of a mean return, the block-bootstrap test of equal Sharpe ratios, the Ledoit–Wolf test of the pooled ratios, and the robustness of the strategy-return tests to the long-run variance estimator.

S13. Decision: future work; keep the ten items in the overview, each marked (not reported), and report them properly in a later draft.
Problem: the overview lists ten items that the appendices do not report. Robustness, statistical (Appendix E): the full model table with the tree ensembles; regime and break tests; remaining variance through the trading day. Robustness, trading (Appendix F): the P&L decomposition; strategy variations; month-ends. Reproducibility (Appendix H): software lock files; hashes of the design matrices; CPU classes; compute time.

B Implementation details

B.1 Samples and windows

- **Estimation and evaluation samples:** Table 4.

E31. Decision: to decide: state the burn-in per panel in the samples note, as 3,125 30-minute grid intervals with its length in trading days, and fix it in make_rehaul_data_methods.py.
Problem: the samples note of Table 4 calls the burn-in 3,125 intervals for every panel, but on the closing-interval series, one row per trading day, it was dropped as 3,125 rows, that is 3,125 trading days, which is why that sample starts in 2010.
On the intraday grid 3,125 intervals are 3,125/48 trading days. The executor now drops the burn-in on the bar grid before the slice (src/backtest/executor.py:137–140, BURNIN_IN_BAR_SPACE, and 933–935), so a rerun of the closing-interval panel would start earlier than the printed sample. The note is generated by make_rehaul_data_methods.py, and re-running it rewrites generated/rehaul_dm_dictionary.tex, which carries the hand edit of note E5, so the note is not regenerated here. No macro holds the burn-in length in trading days.

E20. Decision: two window cells.
Problem: done (2026-10-22, note E34): Panel B prints one window per design, and no printed window is read against fits on another.
Panel B of Table 4 prints $W = 96{,}000$ intervals on the estimation sample of all 48 intervals, the window of the pooled linear models of the model set. Under it, the 13 regular-hours bars carry $W = 13{,}000$ bars (1,000 trading days), the design of the trees and of the pooled linear models of Table 1 and Sections 6–7, and the rolling-forecast row starts for both on the first evaluation day, 3 January 2012. The tuning table (generated/rehaul_dm_tuning.tex, not in the paper) has a row for each design. $W = 24{,}000$ is left only on the row of the tree hyperparameter pre-sample, where it is the window of that search.
Linked note (2026-10-17, user): E24 below, on the number in parentheses in Table 5, is the same underlying window confusion, a printed number read against the wrong estimation window (24,000 intervals, 13,000 bars, 96,000 intervals or 2,000 days).

- **Burn-in:** the first 3,125 samples (30-minute grid intervals), the longest HAR horizon, are dropped (Section 3).
- **Diurnal adjustment** $s^2_{t,i}$ (Section 5.1): $n_{\mathrm{d}} = 20$ trading days, at least $\underline{n}_{\mathrm{d}} = 5$ of them; intervals with fewer have $s^2_{t,i} = 1$ and are not scored.
- **Winsorization** of the target: $n_{\mathrm{w}} = 240$ intervals, all intervals of the day pooled; the bounds are the empirical 1st and 99th percentiles, $\underline{Q}_{t,i} = Q_{0.01}(\mathcal{Y}_{t,i})$ and $\overline{Q}_{t,i} = Q_{0.99}(\mathcal{Y}_{t,i})$.
- **Hyperparameter-selection splits:** $n_{\mathrm{gap}} = 25$ 30-minute intervals and $n_{\mathrm{val}} = 125$ 30-minute intervals.

E21. Decision: label each length with its unit and which step uses it.
Problem: $n_{\text{gap}} = 25$ and $n_{\text{val}} = 125$ are called observations, which on the intraday grid are 30-minute intervals, and the tuning text and Table 4 use the same symbols as numbers of trading days.
This item sets the embargo n_{gap} equal to 25 and the validation length n_{val} equal to 125, and it calls both of them observations. On the intraday grid an observation is one 30-minute interval, so an observation count is 48 times a count of trading days. The linear tuning item, a different step, describes its pre-sample in trading days: four folds of 250 days, running from 6 January 2016 to 2 January 2020. The light re-tune then says it uses two folds of 125 days, which puts the same symbol 125 on a count of days. Table 4 prints that pre-sample as four folds of 250 days after a gap of 25 days. Read as observations, 125 is a number of half hours. Read as in the re-tune and the table, it is a number of trading days.
Decision (2026-10-17, user): the main model is the intraday model, so both lengths are 30-minute intervals: $n_{\text{gap}} = 25$ 30-minute intervals and $n_{\text{val}} = 125$ 30-minute intervals, with no trading-day reference for either. Table 4 and its note, which are generated (`writeup/make_rehaul_data_methods.py`, near line 2283, "-day gap'), still print a 25-day gap.

B.2 Training schemes in practice

- **Closing-interval scheme** ($\mathcal{I} = \{c\}$): Appendix G.3.
- **Intraday scheme** ($\mathcal{I} = \{1, \ldots, M\}$), the paper's training scheme: estimated on all $M = 48$ intervals of each day, overnight included; one observation per interval, so a period is an intraday interval and $\mathcal{W}_{t,i} = \{(t, i-W), \ldots, (t, i-1)\}$ holds the W intervals before (t, i). The full predictor set has 586 predictors (Section 4).
  - For the closing-interval models and the straddle strategy, only the forecasts of the closing interval $i = c$, made at the 15:30 ET forecast origin, are scored and traded; the pooled linear models and tree ensembles of Table 1 are scored on every one of the 13 regular-hours bars.
  - Linear models: $W = 13{,}000$ bars (1,000 trading days of the 13 regular-hours bars), re-estimated every interval.
  - Trees: $W = 13{,}000$ bars (1,000 trading days of the 13 regular-hours bars), the design of the linear models of Table 1; one fixed hyperparameter configuration per tree method, never re-selected, re-estimated every bar.

B.3 Tuned hyperparameters

- **Linear penalties,** chosen by hyperparameter selection on a pre-sample (Section B.4), one grid value per decade (Table 5):

$$\lambda^{\text{ridge}} \in \{10^{-2}, \ldots, 10^{3}\}, \qquad \lambda^{\text{lasso}}, \lambda^{\text{enet}} \in \{10^{-6}, \ldots, 10^{-2}\}.$$

  For lasso and elastic net these are the scikit-learn parameters (Section 5.2).
  - closing-interval linear models (Table 5, Panel A): on the full predictor set ridge 10^{3}, lasso 10^{-2} and elastic net 10^{-2} (top edge); on the HAR benchmark ridge 10^{-2}, lasso 10^{-6} and elastic net 10^{-6} (the lower edge of each grid);
  - pooled linear models of Table 1 and Sections 6–7: one penalty per estimator and predictor set, chosen by the same criterion on 4 validation folds covering 2008–2011; Table 5 does not print these picks (note N24);
  - every penalty is held fixed at every forecast origin of the evaluation span, with no re-tuning (note E23).

N24. Decision: freeze the pick (per E23): each penalty is the single pre-sample value in Panel A, held fixed with no re-tuning; the margin re-tune is not the rule.
Problem: the forecasts are done (2026-10-22). Since 2026-10-18 the closing-interval linear forecasts use the frozen pre-sample picks, with no re-tuning: the master table's subpt_* rows are the frozen arms of the pre-tune study, c_frozen on the HAR benchmark and d_frozen_bb on the full predictor set (experiments/stack_pretune_yhat.py, f_series, and experiments/master_table_close.py, pretune_series; commit 80e7476), which hold the Panel A pick at every forecast origin. results/spxw_pnl/yhat_subpt_*.parquet equal those arms of results/linear_pretune_2026-10-06/forecasts.parquet bit for bit (checked 2026-10-22). Under the margin re-tune, 4 of 24 re-selections had replaced a HAR-benchmark penalty and none a full-set penalty, so only the three HAR-benchmark forecasts changed.
Still open: the Panel A table note still describes the re-tune ("this happened 4 times in 24 re-selections"). It comes from make_rehaul_data_methods.py, which cannot be re-run without wiping the E5 hand edit, so it must be parked or moved to another generator. The picks of the pooled linear models of Table 1 (results/pooled_linear_rerun_2026-10-20_13_2012_w1000/frozen_penalties.json) differ from Panel A for the lasso on both predictor sets, and no table or macro prints them.

E23. Decision: one pre-sample pick per estimator and predictor set, held fixed for the whole evaluation; no causal adaptive re-tuning during evaluation.
Problem: the closing linear forecasts were re-run with re-tuning switched off on 2026-10-18 (note N24), and the pooled linear models of Table 1 hold the penalties chosen on 4 folds covering 2008–2011 (results/pooled_linear_rerun_2026-10-20_13_2012_w1000/frozen_penalties.json). The margin re-tune still enters the six 48-interval pooled linear forecasts of the master table's model set (poolpt_*: arm f of results/pool_linear_pretune_2026-10-06, experiments/stack_pool_pretune_yhat.py, f_series), which count in the aggregates over the 112 forecasts of the model set (Appendices F and G.5) but are not in Table 1. The robustness claim below still needs a sensitivity check to back it.
Each estimator and predictor set gets one penalty, chosen once by hyperparameter selection on the fixed pre-sample window and held constant at every forecast origin. Pre-tuning on a fixed window and holding the penalty constant is a simplification, and it is justified by robustness: the forecasts are not very sensitive to the penalty, so a penalty re-selected during the evaluation would change them little. The text does not claim that the exact value matters. The item above reports the pre-sample pick so that the forecasts can be reproduced, and the decision rests on that insensitivity, which a sensitivity check of the forecasts at neighboring grid values would have to confirm.

- **Trees,** fixed hyperparameters from a search on a 1998–2002 pre-sample (Table 5):
  - LightGBM and XGBoost: the best trial of an Optuna search (Akiba et al., 2019);
  - random forest: trial 255 of 400 ($B = 485$ trees, maximum depth 20, share of predictors tried at each split 0.90, bootstrap sample share 0.41);
  - minimum leaf size $n_{\min}$ (observations per leaf): the number outside the parentheses in Table 5 is the minimum leaf size selected on a window of 24,000 intervals and used, unscaled, by the intraday trees on their 13,000-bar window (Section B.2); the number in parentheses is that minimum leaf size rescaled to the 2,000-day window of the closing-interval trees, $n_{\min} \leftarrow n_{\min} \cdot 2{,}000/24{,}000$, and it is used by the closing-interval trees only.

E24. Decision: the parenthetical is the closing-interval leaf size.
Problem: this sentence calls the rescaled $n_{\min}$ the intraday trees' window. The intraday fits use the unscaled $n_{\min}$, selected on 24,000 intervals. The number in parentheses is the closing-interval leaf size. The leaf minimum $n_{\min}$ was selected on a window of 24,000 intervals. This item rescales it by the ratio of window lengths, $n_{\min} \leftarrow n_{\min} \cdot 2{,}000/24{,}000$, and the sentence calls that product the window of the intraday trees. The note under Table 5 assigns the two numbers the other way around. The figure outside the parentheses is the minimum selected on the 24,000-interval window, and the figure inside the parentheses is that minimum rescaled to a window of 2,000 days. The intraday trees are now the base-5 refit (note E34), estimated on a window of 13,000 regular-hours bars (1,000 trading days; Section B.2), not on the 24,000-interval window of the selection, and they use the outside number as selected, not rescaled to their window. The closing-interval trees are re-estimated every 10 days on a window of 2,000 days, and those are the fits that use the parenthetical.
Linked note (2026-10-17, user): E20 above, on the window cells in Table 4, is the same underlying window confusion, a printed number read against the wrong estimation window (24,000 intervals, 13,000 bars, 96,000 intervals or 2,000 days).
Relabeled (2026-10-17, user): the leaf-minimum item above now calls the number in parentheses the closing-interval trees' minimum leaf size.

B.4 Tuning procedures

- **Linear:**
  - hyperparameter selection on validation MSE over the folds of a pre-sample: for the closing-interval linear models four folds of 250 trading days (6 January 2016 to 2 January 2020), and for the pooled linear models of Table 1 4 folds covering 2008–2011, before the first forecast:
$$\hat{\lambda} = \arg\min_{\lambda} \frac{1}{4} \sum_{v=1}^{4} \mathrm{MSE}_v^{\mathrm{val}}(\lambda),$$
    with $\mathrm{MSE}_v^{\mathrm{val}}(\lambda)$ the mean squared error of $\tilde{y}$ on the validation sample of fold v;
  - $\hat{\lambda}$ is then held fixed at every forecast origin of the evaluation span, with no re-tuning (note E23); the closing-interval forecasts are the frozen arms of the pre-tune study (note N24).
- **Pooled trees** (Table 1 and Sections 6–7): one configuration per method (Section B.3), never re-selected, re-estimated every bar (Section B.2).
- **Closing-interval trees,** in increasing order of tuning effort:
  - fixed hyperparameters, re-estimated every 10 trading days or daily;
  - random search every 250 trading days;
  - Optuna (Akiba et al., 2019): 50 trials, re-tuned every 1, 5, 25 or 250 trading days, best of the first 10, 25 or 50 trials.

E22. Decision: say which rung is in Table 1, and park the others.
Problem: the schemes section, the three-rung list in this subsection, and the draft or preliminary flag on the full-set tree rows each claim a tree cell in Table 1, and no cell is tied to one of them.
A tree cell in Table 1 is the loss of one fitting procedure,and three procedures are written down as if each of them were that procedure. Section B.2 gives each tree method a single configuration, chosen once and never re-selected, and re-estimates it every bar on a window of $W = 13{,}000$ bars. The list in this subsection gives three grades of tuning instead: fixed hyperparameters re-estimated every 10 trading days or every day, a fresh random search every 250 trading days, and an Optuna search of 50 trials re-tuned every 1, 5, 25, or 250 trading days. Section 6.1 adds that the full-set LightGBM and XGBoost rows are an earlier draft and the random forest is preliminary, so those rows need not be any of the procedures just listed. Until one procedure is marked as the source of a given QLIKE cell, the cell has no window, no refit cadence, and no search attached to it.
Update (2026-10-22, third-list): the three-rung list is now headed as the closing-interval trees', and the pooled trees of Table 1 have their own item above (one configuration per method, never re-selected, re-estimated every bar; note E34); Section 6.1 no longer flags the full-set tree rows as an earlier draft or preliminary.

- **Every selection uses only data available at the forecast origin;** no forecast is scored on a day that entered the selection of its own hyperparameters, and the evaluation span $\mathcal{T}_{\text{oos}}$ starts after the pre-sample.

C Data and predictors

- **Sources and spans.** Eight groups, in Section 3. Closing-interval models use all 47 exogenous series as fitted (the 46 of the dictionary and the FOMC statement-interval dummy, dropped on 22 October 2026; note E35). Intraday models use the 46 series in every group, FOMC included.
- **Predictor dictionary.** Appendix Table 6 defines each series. Trailing averages, activity indicators, and calendar columns are added in Section 4.
- **Missing and stale values.** For an exogenous series u, the availability indicator $a_{t,i} = \mathbf{1}\{u \text{ observed at } (t,i)\}$ is recorded before any filling (Section 3). The series is then forward-filled with no staleness limit. On every interval with $a_{t,i} = 0$ the stored regressor is the imputation of Section 4, step 4: the median of the series' observed adjusted values on earlier retained intervals, or the adjusted value on its first observed interval when no earlier observation exists.
- **Duplicate and constant predictors.** At each estimation the window screen of Section 4 drops a column that is constant on $\mathcal{W}_{t,i}$ or an exact bitwise copy of an earlier kept column. The first copy is kept. On the closing-interval model with the full predictor set the screen keeps a median of 369 of the 628 columns.
- **SPXW chain.** Quotes on S&P 500 Weeklys that expire the same day, every 30 minutes: bid, ask, mid-quote, underlying level, implied volatility, and delta. The file covers 1,291 expirations, 2020–2025. The straddle in Section 7.1 uses the 15:30 quotes. A day is dropped when the wider of the two strikes is more than 10 index points from S_t, or when fewer than 10 contracts are quoted at that stamp. S_t is the 15:30 underlying print on the chain, and the official close is the S&P 500 close.
- **Implied-volatility inversion.** The vendor quote on each leg is a volatility per hour, from a geometric bisection. A solve that does not converge returns a bracket node of that bisection; those nodes are set to missing before the two legs are averaged, and a missing leg makes the package's implied variance missing. The 16:00 stamp is not used. The conversion from the hourly quote to the closing half hour is the display in Section 7.1.

- **FOMC calendar.** The four scheduled-announcement series of Section 3, on closing-interval models only. The schedule is known before the interval.

D Hyperparameter selection by model

- **Penalty parameters.** Ridge, lasso, and elastic net minimize the raw sum of squares plus the penalty in Section 5.2. λ is the coefficient on that penalty for ridge. For lasso and elastic net the printed grid is the scikit-learn parameter, and the coefficient in the sum is $2n$ times the printed value. Lasso is $\alpha = 1$. Elastic net is $\alpha = 0.5$. The block weights ω_j are 0.01 on the HAR benchmark, 1 on an additional predictor, and 0 on the intercept. The HAR benchmark fitted alone uses $\omega_j = 1$ on every predictor.
- **Where λ is chosen.** One grid value per decade, Table 5. The pre-sample is four folds of 250 trading days, 6 January 2016 to 2 January 2020, for the closing-interval linear models and 4 validation folds covering 2008–2011 for the pooled linear models of Table 1; the criterion is validation MSE (Section B.4). The value is chosen once on this pre-sample and held fixed for the whole evaluation span, with no re-tuning during evaluation (note E23). Panel A of the table gives the closing-interval picks, the penalty of every closing-interval linear forecast at every forecast origin; on the full predictor set each is at the top of its grid.

  E29. Decision: state the yearly cadence and the closing-interval HAR-benchmark scope, consistent with E23's one pre-sample pick.
  Problem: resolved (2026-10-22, third-list): the light re-tune no longer enters any closing-interval forecast (the frozen arms since 2026-10-18, note N24) and is parked in the tuning procedures (Section B.4), so neither its cadence nor its scope needs stating. For the record: it ran every 250 trading days, on four dates from 23 June 2020 to 16 June 2023, on the closing-interval linear arms only, and all 4 switches were on the HAR-benchmark set (results/linear_pretune_2026-10-06/retune_log.csv; writeup/make_table_close_main.py, about lines 2303–2323). The pooled linear models of Table 1 have frozen penalties (note E23); the six 48-interval pooled linear forecasts of the master table's model set (poolpt_*) still carry their own study's re-tune (note E23).
- **Trees.** LightGBM and XGBoost take the best trial of an Optuna search on a 1998–2002 pre-sample. The random forest takes trial 255 of 400. The selected configuration and the leaf minimum are in Section B.3. The tree cells of Table 1 are the pooled trees, which hold this configuration fixed and are re-estimated every bar; the three tuning rungs in Section B.4 are the closing-interval trees' (note E22).
- **What may enter a choice.** A selection uses only data available at the forecast origin. The evaluation span starts after the pre-sample.

E Robustness of the statistical results

- **Reported comparisons.** Forecast accuracy of the pooled linear models and tree ensembles is Table 1, from 3 January 2012 to 30 April 2024. The Diebold–Mariano statistic, with the Harvey–Leybourne–Newbold correction, the Newey–West variance and the tests of strategy returns are Appendix I.
- **Model confidence set, superior predictive ability, and StepM.** These three procedures are named in that appendix for the comparison of predictor sets. No table reports an MCS membership, an SPA p-value, or a StepM rejection (note E14).
- **Regime and break tests.** Regime tests and break tests have no reported output.
- **Closing-interval against intraday.** Both schemes forecast the same closing half hour. The comparison of their average QLIKE is in Section 6. A forecast of the variance remaining from

Table 2: Sharpe ratio change under alternative definitions of the Sharpe ratio

	All 614 post-warmup days			SR by subsample	
	SR (paper)	Per-year, N_y-weighted	Pooled, $\sqrt{252}$	Before 16 May 2022	From 16 May 2022
Ridge	0.25	−0.22	−0.02	−0.26	0.25
	[0.657]		[0.963]		
Lasso	0.37	−0.26	0.00	−0.20	0.36
	[0.496]		[0.996]		
Elastic net	0.40	−0.17	0.11	−0.14	0.34
	[0.360]		[0.754]		
LightGBM	0.86	0.11	0.55	1.04	0.39
	[0.170]		[0.285]		
XGBoost	0.97	0.31	0.69	1.25	0.51
	[0.180]		[0.181]		
Random forest	0.23	−0.18	0.01	0.30	0.03
	[0.739]		[0.985]		

Notes: Each row is a pooled intraday model class; each entry is the Sharpe ratio of the sign strategy at mid-quotes on the full predictor set minus that on the HAR benchmark. SR is the paper's Sharpe ratio (Section 7.2): the equal-weight mean over calendar years of SR_y, each year annualized by the square root of its option days per calendar year; in brackets, the p-value of the paired block-bootstrap test of equal Sharpe ratios (Section I). Per-year, N_y-weighted: each year's ratio scaled by the square root of its number of option days N_y, the years weighted by N_y. Pooled, $\sqrt{252}$: the ratio of the mean to the standard deviation of all 614 daily returns, times $\sqrt{252}$; in brackets, the p-value of the Ledoit and Wolf (2008) test, which tests the equality of these pooled ratios. The last two columns give SR on the 125 post-warmup days before 16 May 2022, the first week with an expiration on every weekday, and on the 489 days from that date, each subsample annualized as a sample of its own. *, ** and *** denote significance at the 10%, 5% and 1% levels.

an earlier origin to the close is not reported.

F Robustness of the trade

- **Definition of the Sharpe ratio.** Table 2 repeats the comparison of the full predictor set with the HAR benchmark for the 6 pooled intraday model classes under the Sharpe ratio SR of Section 7.2 and under two other definitions (note D6). Under SR the change is positive in 6 of the 6 classes (from +0.23 to +0.97), and the bootstrap test finds it statistically significant at the 5% level in 0. With each year scaled by $\sqrt{N_y}$ and the years weighted by N_y, the change is positive in 2 classes (from −0.26 to +0.31). For the ratio of the pooled daily returns scaled by $\sqrt{252}$, it is positive in 4 classes (from −0.02 to +0.69), and the Ledoit and Wolf (2008) test finds it statistically significant at the 5% level in 0. Computing SR separately on the 125 post-warmup days before 16 May 2022 and on the 489 days from that date, the change is positive in 3 and 6 classes, respectively (from −0.26 to +1.25 and from +0.03 to +0.51).
- **Date-shift test.** The sign strategy is re-scored with the forecast of the interval j half hours away from the closing interval, placed on the closing interval's diurnal adjustment:

$$q_t = \mathrm{sign}(\hat{y}^2_{t,c+j}\, s^2_{t,c} - \mathrm{IV}^2_t).$$

$j < 0$ reads an earlier interval of day t. $j > 0$ reads day $t + 1$. The test is scored on the 9

forecasts whose output holds every intraday interval, the earlier draft's (note N18), on the 865 option-data days on which every shifted forecast exists, with the Sharpe ratio annualized by one fixed factor on all those days rather than year by year on the 614 post-warmup days of Section 7.2 (note N15b). Across j, the median of that Sharpe ratio runs from 0.00 to 1.58 (a swing of ± 0.79), inside the ± 1.06 95% sampling band of a Sharpe ratio so annualized under i.i.d. returns.

N18. Decision: to decide (data export needed).
Problem: the shift test runs on the earlier draft's nine pooled 48-bar forecasts (a0, blk2, blk2_inc, lgbm, xgb, lasso_t, lasso_f, enet, b2lasso); the lgbm and xgb rows there are the earlier draft's trees, not the pooled trees of this paper. Of the paper's 12 pooled intraday forecasts, the six pooled trees (the base-5 refit of note E34, results/pooled_trees_rerun13_2026-10-20/w1000) have every one of the 13 regular-hours bars, so they can be shifted for $j < 0$ within the regular-hours day; the six pooled linears are now the 13-bar rerun on the trees' design (note E34; results/pooled_linear_rerun_2026-10-20_13_2012_w1000/forecasts.parquet, read by writeup/pooled_linear_13.py), which also holds every regular-hours bar (the master table's 48-interval poolpt_* rows, results/pool_linear_pretune_2026-10-06/forecasts.parquet, hold only the 16:00 row). Restating the test on the pooled forecasts needs loaders for the two forecasts.parquet tables (2026-10-22, E34: the all-interval export of the pooled linears is no longer needed) in experiments/close_shift_test_headline.py (writeup/pooled_trees_13.py reads only its 16:00 rows; the test's $k = 0$ checks also need a reference for that table). See E26.

- **Murphy diagram.** The elementary score at the implied variance is $\mathrm{ES}_{1,t}$ in Section 7.2. Its rank correlation with the Sharpe ratio, across the 12 pooled intraday forecasts, is in Section 7.3. The subsample split is around May 2022.
- **Profit and loss.** The reported return is R_t in Section 7.1. A decomposition of that return into separate variance and direction components is not reported.
- **Compounding.** Wealth starts at 1 and the premium each day is the fraction $f = 3\%$ in Section 7.2. Across the 112 forecasts, on the 614 post-warmup days, the median terminal value is 2.4 at mid-quotes and 1.4 at quoted bid and ask prices, and the maximum drawdowns reach 59% and 68%.
- **Position sizing.** The sign strategy and the two sized rules, the robust clip and the percentile rank, are in Section 7.1. All three are scored on the 614 days after the warmup of 252 days.
- **Other origins.** No other forecast origin is reported. The interval-shift test above keeps the 15:30 close, implied volatility and straddle, and takes only the model's forecast from another interval. A separate month-end strategy is not reported.

G Closing-interval models

G.1 Closing-interval sample

- The closing-interval specifications incorporate all 47 exogenous series as fitted, the FOMC group included (the dictionary lists 46 series and 3 FOMC series since the statement-interval dummy was dropped on 22 October 2026; note E35); so do the pooled intraday specifications, with 46 series (Section 3).
- The closing-interval sample ($i = c$) consists of 3,469 trading days from July 2010 through April 2024, estimated on a rolling window of $W = 2{,}000$ trading days.
- This rolling procedure yields 1,469 out-of-sample closing-interval forecasts starting in June 2018:

$$1{,}469 = 3{,}469 - 2{,}000.$$

- Hyperparameter selection for the linear specifications is conducted on an initial pre-sample us-

ing closing-interval observations starting on 13 April 2004, with validation windows spanning 6 January 2016 to 2 January 2020.

G.2 Closing-interval predictors

- The horizon set of the closing-interval model is base 5,

$$\mathcal{K} = \{1, 5, 25, 125, 625, 3125\},$$

the same set as the intraday model. Here, k counts retained intervals across the full intraday grid $\mathcal{G}$ (including overnight sessions), while the closing-interval specification restricts estimation to the rows $i = c$.
- In the closing-interval panel, the k-interval trailing averages of the activity indicator, $\bar{\chi}_{t,i}^{(k)}$, are computed for the 6 mode-inflated series of the dictionary and, in the fitted forecasts, for the FOMC statement-interval dummy dropped on 22 October 2026 (notes E6, E35).
- Under the closing-interval specification with the full predictor set, the window filter retains a median of 369 out of 628 candidate columns across rolling windows.
- **Predictor set dimensions.** In the closing-interval specification, the HAR benchmark contains 22 candidate columns, whereas the full predictor set contains 628 columns:

$$628 = \underbrace{22}_{\bar{y}^{(k)},\ \boldsymbol{D}} + \underbrace{47 \times 6}_{\bar{u}^{(k)}} + \underbrace{47 \times 6}_{\bar{a}^{(k)}} + \underbrace{(6+1) \times 6}_{\bar{\chi}^{(k)}}, \qquad 22 = 6 + 16.$$

G.3 Closing-interval HAR benchmark and estimation

- **Closing-interval HAR benchmark.** The benchmark forecast of the main text is the daily HAR regression of Corsi (2009), estimated by OLS on the HAR components and calendar dummies (Section 5.2) using closing-interval observations only: one observation per trading day, re-estimated daily on a rolling window of $W = 2{,}000$ trading days.
- Under the closing-interval scheme, a period is a trading day and the target is $y_t \equiv y_{t,c}$; the refit cadence of each model class is in the next item.
- **Closing-interval scheme:** $W = 2{,}000$ trading days for every model class; re-estimated every trading day, except the tree specifications re-estimated every 10 trading days (fixed hyperparameters or random search).

E32. Decision: user decision 2026-10-18: the tree specifications refit every session (every trading session), not every 10 trading days; "daily" applies only to the closing-interval linear models. The rerun is still pending.
Problem: the closing-interval tree specifications refit every 10 trading days (specs/causal_tune_trees.py:35–40, REFIT_EVERY = 10 sessions; jobs/sge/unification_tree.sge:17, TREE_REFIT_EVERY=10), so their printed numbers are not every-session refits.
Requires a rerun of the closing-interval tree specifications that refit every 10 trading days, with an every-session refit (the fixed-hyperparameter rung, subtree_daily_*, already refits every session: experiments/close_studies_export.py:75 and 80, refit_every = 1). The pooled intraday models already refit at every interval (experiments/pool_linear_pretune.py:8–9 and specs/causal_tune_linear.py:159, REFIT_FREQUENCY = 1; experiments/pooled_trees_score.py:139–140), and the closing-interval linear models and the HAR benchmark every trading day (experiments/close_linear_pretune.py:11).

G.4 Closing-interval standardization

The closing-interval model of Section G.5 standardizes the averages of series values on its own window $\mathcal{W}_t$ (Section B.2). Rows t' are trading days ($i = c$ suppressed):

$$z_t^{(k)} = \frac{\bar{u}_t^{(k)} - \mathrm{med}_{t' \in \mathcal{W}_t}\, \bar{u}_{t'}^{(k)}}{\max\{\mathrm{IQR}_t^{(k)},\ 0.01 \max_{t' \le t} \mathrm{IQR}_{t'}^{(k)},\ \mathrm{IQR}_t^{*(k)}\}}, \qquad \mathrm{IQR}_t^{(k)} = \mathrm{IQR}\{\bar{u}_{t'}^{(k)} : t' \in \mathcal{W}_t\}.$$

- $\mathrm{IQR}_t^{*(k)}$: IQR of $\bar{u}_{t'}^{(k)}$, $t' < t$, over rows whose k-interval averaging window is fully observed.
- Column whose mode accounts for at least 20% of rows: $\mathrm{IQR}_t^{*(k)}$ also excludes rows at the mode.
- $\mathrm{IQR}_t^{*(k)} = 1$ until 8 such rows exist.
- $\bar{a}^{(k)}$, $\bar{\chi}^{(k)}$: not standardized. HAR components, calendar dummies: $\mathrm{IQR}_t^{*(k)} = 0$.
- $t < W$: the first window's statistics.

G.5 Training on the last 30 minutes

- **Closing-interval estimation scheme ($\mathcal{I} = \{c\}$).** As a final refinement, we restrict estimation of each model class in Section 5.2 exclusively to observations from the closing half hour $i = c$ (15:30–16:00 ET), matching the estimation domain of the daily HAR benchmark.
  - Under this scheme, each trading day contributes a single observation, so a period corresponds to one trading day and the rolling window $\mathcal{W}_t = \{t - W, \ldots, t - 1\}$ spans $W = 2{,}000$ trading days (Appendix B.2). Rolling standardization of the exogenous series averages on $\mathcal{W}_t$ follows Appendix G.4.
  - We hold the predictor sets, evaluation sample $\mathcal{T}_{\mathrm{oos}}$, and straddle trading rules unchanged; statistical accuracy is reported in Table 3 for the 15:30 forecasts of the closing half hour over $\mathcal{T}_{\mathrm{oos}}$.
- **Improvements in QLIKE accuracy.** Training solely on the closing interval reduces mean QLIKE relative to pooled intraday training in every model class and under both predictor sets. Specifically, the percentage excess loss of the pooled intraday forecast $f_{m,\mathcal{X},\{1,\ldots,M\}}$ relative to the closing-interval forecast $f_{m,\mathcal{X},\{c\}}$, $\Delta\% = 100(\bar{L}^{\mathrm{pool}}/\bar{L}^{\mathrm{close}} - 1)$, is strictly positive throughout:
  - for penalized linear models, pooling increases mean QLIKE by +16.5% to +16.7% on the HAR benchmark (DM from 2.97 to 2.99, all significant at the 5% level) and by +17.7% to +24.0% on the full predictor set (significant for the lasso);
  - for tree ensembles, pooling increases mean QLIKE by +4.6% to +11.4% on the HAR benchmark and by +7.3% to +8.0% on the full predictor set; the difference is significant at the 5% level for 1 of the three tree classes on the HAR benchmark and for 2 on the full predictor set.
- **Preservation of the exogenous-predictor gains.** Within the closing-interval scheme, expanding from the HAR benchmark to the full predictor set lowers mean QLIKE in all 6 model classes, with relative reductions $\Delta\%$ between -4.1% and -0.1%, although none of the single-class reductions is individually significant at the 5% level (for lasso, $\Delta\% = -4.0\%$, $\mathrm{DM} = -1.91$, $p = 0.056$).
  - Across all 51 matched predictor-set pairs in the full model set (45 tree and 3 linear specifications trained on the closing interval, plus 3 pooled intraday specifications), adding

the exogenous predictors reduces mean QLIKE in 35 pairs, 4 of which are significant at the 5% level.

- Across the 12 pooled intraday forecasts, 0 achieve a lower mean QLIKE than the daily HAR benchmark (0 significantly), while 10 perform significantly worse; the closing-interval forecasts are left out of this count.
- Comparing tree ensembles directly against the closing-interval ridge regression on the same predictor set, the tree specification incurs a higher mean QLIKE in 48 of 48 pairs (23 significantly).
- **Sensitivity of squared-error loss to extreme days.** In contrast to QLIKE, evaluated under MSE the full predictor set increases squared-error loss in 49 of the 51 matched pairs across the model set (0 significantly).
  - A single date outside the option-data subsample, 19 March 2020, accounts for 28–39% of the total cumulative squared error of the closing-interval tree models over $\mathcal{T}_{\text{oos}}$, compared with at most 0.9% of their total QLIKE loss.
  - Similarly, five trading days in March 2020 account for 73–87% of the cumulative squared error of the daily HAR benchmark and of the closing-interval lasso and XGBoost on the full predictor set; when February–April 2020 are excluded, the closing-interval lasso achieves a lower MSE than the benchmark forecast.
- **Economic value under the closing-interval scheme** (sign strategy evaluated on the 614 post-warmup days of $\mathcal{T}_{\text{opt}}$):
  - The closing-interval elastic net on the full predictor set attains SR = 1.68 at mid-quotes and 1.33 at quoted bid/ask prices; its mean return gain over the short-straddle benchmark has $t = 1.92$ (fixed-b $p = 0.072$). The closing-interval lasso on the full predictor set attains SR = 1.44 at mid-quotes and 1.09 at quoted bid/ask prices.
  - Closing-interval LightGBM on the full predictor set, with fixed hyperparameters re-estimated every trading day (the specification of the closing-interval tree rows of the main results table), attains SR = 1.19 at mid-quotes; the tree rows of the straddle table are the pooled intraday trees.
  - Replacing the HAR benchmark with the full predictor set changes the closing-interval sign-strategy Sharpe ratio by $\Delta\text{SR} = +0.21$ for the elastic net and -0.03 for the lasso (neither statistically significant at the 5% level by the block-bootstrap test of equal Sharpe ratios); across all 51 matched predictor-set pairs in the model set, $\Delta\text{SR} > 0$ in 47 pairs (median +0.71), and the same test finds a statistically significant gain at the 5% level in 13 pairs and a statistically significant loss in 0.
  - Relative to the unconditional short-straddle benchmark, $\Delta\text{SR} > 0$ for 106 of the 112 forecasts in the model set; the mean return gain is statistically significant at the 5% level (fixed-b test) for 0 of them, 0 of which use the full predictor set.
  - Finally, none of the mean return differentials between a pooled intraday model and the corresponding closing-interval model of the same class is statistically significant at the 5% level (fixed-b test of the mean daily return differential, bandwidth $S = 33$, on the 614 post-warmup days).

H Reproducibility

- **Generator of each reported number.**

Table 3: Forecast accuracy when training on the last 30 minutes

	HAR benchmark			All features		
Model	MSE (10^{-10})	QLIKE	R^2_{OOS} (%)	MSE (10^{-10})	QLIKE	R^2_{OOS} (%)
OLS (benchmark forecast)	4.15	0.1033	78.2	–	–	–
Panel A: Linear models						
Ridge	4.15	0.1033	78.2	7.30	0.1031	61.7
Lasso	4.15	0.1033	78.2	5.94	0.0992	68.8
Elastic net	4.15	0.1033	78.2	6.89	0.1014	63.9
Panel B: Tree ensembles						
LightGBM	5.75	0.1094	69.8	17.19	0.1049	9.8
XGBoost	5.27	0.1076	72.4	20.94	0.1031	−9.9
Random forest	4.62	0.1068	75.8	14.15	0.1060	25.8

Notes: Rolling forecasts from 3 January 2020 to 30 April 2024, every forecast made with a model estimated on its trailing window only, so every forecast is out of sample; each is scored against the actual realized variance of the closing half hour, RV_t, on the $T_{\text{oos}} = 1{,}082$ trading days on which every forecast of both training schemes exists, early closes dropped (Section 5.3). The span starts with the intraday linear forecasts, on the first trading day after the pre-sample on which the penalties were selected (6 January 2016 to 2 January 2020). The MSE entries are in units of 10^{-10}: the printed number times 10^{-10} is the mean squared error $T_{\text{oos}}^{-1}\sum_t(\text{RV}_t - \hat{\sigma}_t^2)^2$. QLIKE is the mean QLIKE loss. $R^2_{\text{OOS}} = 1 - \sum_t(\text{RV}_t - \hat{\sigma}_t^2)^2 / \sum_t(\text{RV}_t - \overline{\text{RV}}_{t-1})^2$, in percent, where the benchmark $\overline{\text{RV}}_{t-1}$ is the historical mean of the closing-half-hour realized variance over all trading days from 5 January 1998 to day $t-1$ (Campbell and Thompson, 2008). All features is the full predictor set of Section 4; the OLS is estimated on the HAR benchmark only (–: no forecast). Each model is estimated on the last 30 minutes of each trading day only (one observation per day); the OLS on the HAR benchmark is the benchmark forecast. Trees use fixed hyperparameters and are re-estimated daily (Section B.2). The intraday models on the same days are in Table 1.

- `writeup/make_rehaul_data_methods.py`: samples, windows, hyperparameter tables, and the predictor dictionary.
- `writeup/make_rehaul_oos_rerun.py`: Table 1.
- `writeup/make_rehaul_oos_table.py`: Table 3.
- `writeup/make_rehaul_results.py`: the accuracy macros in Section 6.
- `writeup/make_rehaul_economics.py`: the straddle macros in Section 7.
- `writeup/make_rehaul_fi.py`: the penalty weight on the HAR block in Appendix D and the closing-interval forecast count in Appendix G.1.
- `writeup/make_rehaul_appendix.py`: the appendix-table macros.
- `writeup/make_table_close_main.py`: the closing-interval master-table macros, including the closing-interval predictor counts.
- `writeup/make_fig_qlike_sharpe.py`: the matched predictor-set pair counts.
- `writeup/hac_inference.py`: the fixed-b test of a mean return, the Sharpe ratio of Section 7.2 and the block-bootstrap test of its differences, and the Ledoit–Wolf test of equal Sharpe ratios, used by the three generators above; `writeup/make_rehaul_results.py` writes Table 7 and `writeup/make_rehaul_economics.py` Table 2.
- `writeup/pooled_trees_13.py`: the 16:00 forecasts of the pooled tree ensembles (results/pooled_trees_rerun13_2026-10-20/w1000, fitted by `experiments/pooled_trees_rerun_13.p` read by `writeup/make_rehaul_results.py`, `writeup/make_rehaul_economics.py` and `writeup/make_rehaul_oos_table.py`.
- `writeup/pooled_linear_13.py`: the 16:00 forecasts of the pooled linear mod-

els of Table 1 (results/pooled_linear_rerun_2026-10-20_13_2012_w1000, fitted by `experiments/pooled_linear_rerun_2026_10_19_13.py`), read by `writeup/make_rehaul_results.py` and `writeup/make_rehaul_economics.py`.
  - `writeup/make_buckets_tex.py`: the group sizes and the exogenous series count.
  - `experiments/close_murphy_splits_headline.py`: the Murphy-diagram macros.
  - `experiments/close_compounding_headline.py`: the compounding macros in Appendix F.
  - `experiments/close_shift_test_headline.py`: the interval-shift test macros in Appendix F.
- **Software list.** `requirements.txt` at the repository root names the packages and a version range for each. It is not a pinned lock file, and the draft does not record a second lock file.
- **Design matrices, CPU class, and compute time.** The generators read the forecasts and tables already on disk. The draft does not record a hash of the design matrix behind Table 1, the CPU class of those fits, or their compute time.

Table 4: Estimation and evaluation samples.

Sample	Start	End	Obs.	Window W
Intraday grid: 30-minute intervals with an observed $RV_{t,i}$, outside the weekend break	1998-01-05	2024-04-30	303,442	–
Panel A: Closing-interval models (trading days)				
Closing-interval observations, after the burn-in	2010-07-12	2024-04-30	3,469	2,000
Rolling-window forecasts	2018-06-25	2024-04-30	1,469	–
Hyperparameter-selection pre-sample (linear models)	2004-04-13	2020-01-02	3,947	2,000
Validation samples (4×250 days)	2016-01-06	2020-01-02	1,000	2,000
Panel B: Intraday models (30-minute intervals)				
Estimation sample, after the 3,125-interval burn-in	1998-06-04	2024-04-30	300,317	96,000
Regular-hours bars, bar ends 10:00–16:00 ET: the trees and the pooled linear models of Table 1 and Sections 6–7	1998-06-04	2024-04-30	85,314	13,000
Rolling-window forecasts from the first evaluation day, less early-close days (26)	2012-01-03	2024-04-30	40,466	–
Tree hyperparameter pre-sample (before 2003)	1998-06-04	2002-12-31	36,285	24,000
Panel C: Evaluation samples (trading days)				
Evaluation span (T_{oos}): trading days on which every forecast exists, less early-close days (6)	2020-01-03	2024-04-30	1,082	–
Option-data subsample (T_{opt}): SPXW expiration days (872) less early-close days (6)	2020-01-03	2024-04-30	866	–
Later SPXW expiration days, held out (413 full trading days)	2024-05-01	2025-12-31	419	–

Notes: Dates are interval-end dates (ET); Obs. and the rolling estimation window W are in the units of each panel. The burn-in, 3,125 intervals, is the longest HAR horizon and is applied to each series as built (the closing-interval series, the intraday grid and, for the linear pre-sample, the series of the two intervals of the last hour), so the closing-interval sample starts in 2010 and the pre-sample in 2004. The trees and the pooled linear models of Table 1 and Sections 6–7 are re-estimated at every bar on the previous 13,000 regular-hours bars (1,000 trading days of 13 bars); the pooled linear models on all 48 intervals (in the model set) use $W = 96{,}000$ intervals (2,000 trading days). The tree hyperparameters are chosen once, by an Optuna search on the 1998–2002 development sample that fits each candidate on the previous 24,000 intervals and scores it on the next 250, stepping forward 250 intervals at a time. Each validation sample follows a 25-day gap and a 2,000-day estimation sample (Section B.1). Full trading days: expiration days that are not early-close days. Panel C: the evaluation span holds every trading day on which all rolling forecasts exist (Section 5.3); the straddle economics use its option-data subsample (Section 7).

Table 5: Hyperparameters: penalized regressions and tree ensembles.

Model	Hyperparameter	Search space	Value: HAR benchmark	Value: Full predictor set
Panel A: Penalized linear models, selected on the pre-sample				
OLS	–	–	–	–
Ridge	Penalty, λ	$\{10^{-2}, 10^{-1}, \ldots, 10^{3}\}$ (6)	10^{-2}	10^{3}
Lasso	Penalty, λ	$\{10^{-6}, 10^{-5}, \ldots, 10^{-2}\}$ (5)	10^{-6}	10^{-2}
Elastic net	Penalty, λ	$\{10^{-6}, 10^{-5}, \ldots, 10^{-2}\}$ (5)	10^{-6}	10^{-2}
	Mixing weight, α	fixed	0.5	0.5
All three	Penalty weight on the HAR benchmark, $\underline{\omega}$	fixed	1	0.01
Panel B: Tree ensembles, fixed hyperparameters (selected on the 1998–2002 pre-sample)				
LightGBM	Leaves per tree, J	[15, 255] (log)	20	
	Learning rate, η	[0.01, 0.3] (log)	0.013	
	Boosting rounds, B	[50, 500]	392	
	Minimum leaf size, $n_{\min}$	[20, 500] (log)	98 (8)	
	Share of predictors per tree	[0.5, 1]	0.70	
	Share of observations per tree	[0.5, 1]	0.56	
	L_1 penalty on leaf weights, λ_1^{tree}	$[10^{-8}, 10]$ (log)	10^{-8}	
	L_2 penalty on leaf weights, λ_2^{tree}	$[10^{-8}, 10]$ (log)	0.016	
XGBoost	Maximum depth	[3, 10]	5	
	Learning rate, η	[0.01, 0.3] (log)	0.0123	
	Boosting rounds, B	[50, 500]	367	
	Minimum leaf size, $n_{\min}$	[1, 100] (log)	50.2 (4.18)	
	Share of observations per tree	[0.5, 1]	0.66	
	Share of predictors per tree	[0.5, 1]	0.65	
	L_1 penalty on leaf weights, λ_1^{tree}	$[10^{-8}, 10]$ (log)	0.0215	
	L_2 penalty on leaf weights, λ_2^{tree}	$[10^{-8}, 10]$ (log)	1.2×10^{-5}	
	Penalty per leaf, γ	$[10^{-8}, 1]$ (log)	1.5×10^{-6}	
Random forest	Number of trees, N	[50, 500]	485	
	Maximum depth	[2, 32] (log)	20	
	Share of predictors per split	[0.02, 1] (log)	0.90	
	Bootstrap sample share	[0.05, 1] (log)	0.41	
	Minimum leaf size, $n_{\min}$	[1, 1000] (log)	12 (1)	

Notes: Panel A: Section 5.2 states the scale of λ. One grid value per decade (grid size in parentheses), selected by mean validation MSE over the 4 folds of the pre-sample (Table 4). The selected penalty is then held fixed at every forecast origin, with no re-tuning. Panel B: the best trial of an Optuna search on the intraday pre-sample (24,000-interval window; random forest: trial 255 of 400), ranked by MSE and then held fixed for both predictor sets; (log): searched on a logarithmic scale. $n_{\min}$ was selected on 24,000 intervals and is rescaled to the 2,000-day window (rescaled value in parentheses). Implementations: LightGBM (Ke et al., 2017) with bagging at every round, XGBoost (Chen and Guestrin, 2016) with histogram-based split finding and the scikit-learn random forest with bootstrap samples; random seeds are fixed.

Table 6: Predictor definitions.

Variable	Description	Source	Diurnal adjustment	Transform $g(v)$	Predictors, $k \in \mathcal{K}$
Realized moments (7; front-month E-mini, one-minute log-price changes)					
`sumret`	Sum of one-minute changes in the log price of the front-month contract, $\texttt{sumret}_{t,i} = \sum_k r_{t,i,k}$	E-mini S&P 500 futures	s.d.	v	$u^{(k)}$, $\bar{a}^{(k)}$
`sumabsret`	Realized absolute variation, $\sum_k \lvert r_{t,i,k}\rvert$	E-mini S&P 500 futures	mean	v	$u^{(k)}$, $\bar{a}^{(k)}$
`sumret3`	Realized skewness, unnormalized, $\sum_k r_{t,i,k}^3$	E-mini S&P 500 futures	s.d.	$v^{1/3}$	$u^{(k)}$, $\bar{a}^{(k)}$
`sumret4`	Realized quarticity, unscaled, $\sum_k r_{t,i,k}^4$	E-mini S&P 500 futures	mean	$\text{sign}(v)\,\lvert v\rvert^{1/4}$	$u^{(k)}$, $\bar{a}^{(k)}$
`sumpret2`	Positive realized semivariance, $\sum_k r_{t,i,k}^2\,\mathbf{1}\{r_{t,i,k} > 0\}$	E-mini S&P 500 futures	mean	$\sqrt{v}$	$u^{(k)}$, $\bar{a}^{(k)}$
`sumbipow`	Bipower variation, $\texttt{sumbipow}_{t,i} = \sum_{k\geq 2} \lvert r_{t,i,k}\rvert\,\lvert r_{t,i,k-1}\rvert$	E-mini S&P 500 futures	mean	$\sqrt{v}$	$u^{(k)}$, $\bar{a}^{(k)}$
`sumautocov`	First-order realized autocovariance, $\texttt{sumautocov}_{t,i} = \sum_{k\geq 2} r_{t,i,k}\,r_{t,i,k-1}$	E-mini S&P 500 futures	s.d.	$\text{sign}(v)\,\lvert v\rvert^{1/2}$	$u^{(k)}$, $\bar{a}^{(k)}$
Liquidity (14)					
`sumvolume`	Number of front-month contracts traded in the half hour	E-mini S&P 500 futures	mean	$\log v$	$u^{(k)}$, $\bar{a}^{(k)}$
`numobs`	Number of non-missing one-minute intervals in the half hour, $\texttt{numobs}_{t,i} \leq 30$; full-sample mode 30 on 72% of intervals	E-mini S&P 500 futures	mean	$\log v$	$u^{(k)}$, $\bar{a}^{(k)}$, $\bar{\chi}^{(k)\dagger}$
`turnover_ewstock`	Equal-weighted turnover, shares traded in the half hour divided by the previous day's shares outstanding	Constituents	mean	$\sqrt{v}$	$u^{(k)}$, $\bar{a}^{(k)}$
`buyturnover_ewstock`	Equal-weighted buyer-initiated turnover, the same ratio using estimated buy orders	Constituents	mean	$\sqrt{v}$	$u^{(k)}$, $\bar{a}^{(k)}$
`sellturnover_ewstock`	Equal-weighted seller-initiated turnover, the same ratio using estimated sell orders	Constituents	mean	$\sqrt{v}$	$u^{(k)}$, $\bar{a}^{(k)}$
`effspread_ewstock`	Equal-weighted effective spread, the mean percent difference between the trade price and the preceding midquote	Constituents	mean	$\sqrt{v}$	$u^{(k)}$, $\bar{a}^{(k)}$
`spread_ewstock`	Equal-weighted quoted spread, the mean percent difference between the bid and the ask	Constituents	mean	$\log v$	$u^{(k)}$, $\bar{a}^{(k)}$

continued on the next page

Table 6 (continued)

Variable	Description	Source	Diurnal adjustment	Transform $g(v)$	Predictors, $k \in \mathcal{K}$
`turnover_vwstock`	Value-weighted turnover, shares traded in the half hour divided by the previous day's shares outstanding	Constituents	mean	$\sqrt{v}$	$u^{(k)}$, $\bar{a}^{(k)}$
`buyturnover_vwstock`	Value-weighted buyer-initiated turnover, the same ratio using estimated buy orders	Constituents	mean	$\sqrt{v}$	$u^{(k)}$, $\bar{a}^{(k)}$
`sellturnover_vwstock`	Value-weighted seller-initiated turnover, the same ratio using estimated sell orders	Constituents	mean	$\sqrt{v}$	$u^{(k)}$, $\bar{a}^{(k)}$
`effspread_vwstock`	Value-weighted effective spread, the mean percent difference between the trade price and the preceding midquote	Constituents	mean	$\sqrt{v}$	$u^{(k)}$, $\bar{a}^{(k)}$
`spread_vwstock`	Value-weighted quoted spread, the mean percent difference between the bid and the ask	Constituents	mean	$\log v$	$u^{(k)}$, $\bar{a}^{(k)}$
`ofi_ewstock`	Equal-weighted order-flow imbalance, (Buy − Sell)/(Buy + Sell) of turnover	Constituents	s.d.	v	$u^{(k)}$, $\bar{a}^{(k)}$
`ofi_vwstock`	Value-weighted order-flow imbalance, (Buy − Sell)/(Buy + Sell) of turnover	Constituents	s.d.	v	$u^{(k)}$, $\bar{a}^{(k)}$
Equal-weighted market (6; simple averages of S&P 500 constituents)					
`sumret2_ewstock`	Equal-weighted realized variance of 5-minute log returns	Constituents	mean	$\sqrt{v}$	$u^{(k)}$, $\bar{a}^{(k)}$
`sumret3_ewstock`	Equal-weighted realized skewness of 5-minute log returns, unnormalized	Constituents	s.d.	$v^{1/3}$	$u^{(k)}$, $\bar{a}^{(k)}$
`sumret4_ewstock`	Equal-weighted realized quarticity of 5-minute log returns, unscaled	Constituents	mean	$\text{sign}(v)\,\lvert v\rvert^{1/4}$	$u^{(k)}$, $\bar{a}^{(k)}$
`sumabsret_ewstock`	Equal-weighted realized absolute variation of 5-minute log returns	Constituents	mean	v	$u^{(k)}$, $\bar{a}^{(k)}$
`sumbipow_ewstock`	Equal-weighted bipower variation of 5-minute log returns	Constituents	mean	$\sqrt{v}$	$u^{(k)}$, $\bar{a}^{(k)}$
`sumpret2_ewstock`	Equal-weighted positive realized semivariance of 5-minute log returns	Constituents	mean	$\sqrt{v}$	$u^{(k)}$, $\bar{a}^{(k)}$
Value-weighted market (6; previous day's market capitalization)					
`sumret2_vwstock`	Value-weighted realized variance of 5-minute log returns	Constituents	mean	$\sqrt{v}$	$u^{(k)}$, $\bar{a}^{(k)}$
`sumret3_vwstock`	Value-weighted realized skewness of 5-minute log returns, unnormalized	Constituents	s.d.	$v^{1/3}$	$u^{(k)}$, $\bar{a}^{(k)}$

continued on the next page

Table 6 (continued)

Variable	Description	Source	Diurnal adjustment	Transform $g(v)$	Predictors, $k \in \mathcal{K}$
`sumret4_vwstock`	Value-weighted realized quarticity of 5-minute log returns, unscaled	Constituents	mean	$\text{sign}(v)\,\lvert v\rvert^{1/4}$	$u^{(k)}$, $\bar{a}^{(k)}$
`sumabsret_vwstock`	Value-weighted realized absolute variation of 5-minute log returns	Constituents	mean	v	$u^{(k)}$, $\bar{a}^{(k)}$
`sumbipow_vwstock`	Value-weighted bipower variation of 5-minute log returns	Constituents	mean	$\sqrt{v}$	$u^{(k)}$, $\bar{a}^{(k)}$
`sumpret2_vwstock`	Value-weighted positive realized semivariance of 5-minute log returns	Constituents	mean	$\sqrt{v}$	$u^{(k)}$, $\bar{a}^{(k)}$
Sentiment (3)					
`stocktwits_attention`	Number of StockTwits posts in the half hour	StockTwits	mean	$\log v$	$u^{(k)}$, $\bar{a}^{(k)}$
`stocktwits_sentiment`	Score of tagged posts, 1 if all are bullish and −1 if all are bearish	StockTwits	none	v	$u^{(k)}$, $\bar{a}^{(k)}$
`stocktwits_sentcount`	Number of posts in the half hour that carry a bullish or bearish tag	StockTwits	mean	$\log v$	$u^{(k)}$, $\bar{a}^{(k)}$
Implied volatility (3)					
`vix`	VIX index (30-day S&P 500 implied volatility)	Cboe	none	$\log v$	$u^{(k)}$, $\bar{a}^{(k)}$
`vvix`	VVIX index (implied volatility of VIX)	Cboe	none	$\log v$	$u^{(k)}$, $\bar{a}^{(k)}$
`vix3m`	VIX3M index (3-month S&P 500 implied volatility)	Cboe	none	$\log v$	$u^{(k)}$, $\bar{a}^{(k)}$
Volatility demand (4; net vega bought by non-market-makers)					
`voldemand_spx_open_and_close`	Net vega bought by non-market-makers, S&P 500 index and top S&P 500 ETF options, trades that open or close a position; full-sample mode 0 on 59% of intervals	Cboe	none	$\text{asinh}(u/\kappa)$	$u^{(k)}$, $\bar{a}^{(k)}$, $\bar{\chi}^{(k)\dagger}$
`voldemand_spx_open_only`	Net vega bought by non-market-makers, S&P 500 index and top S&P 500 ETF options, trades that open a position; full-sample mode 0 on 60% of intervals	Cboe	none	$\text{asinh}(u/\kappa)$	$u^{(k)}$, $\bar{a}^{(k)}$, $\bar{\chi}^{(k)\dagger}$
`voldemand_all_open_and_close`	Net vega bought by non-market-makers, all traded options, trades that open or close a position; full-sample mode 0 on 59% of intervals	Cboe	none	$\text{asinh}(u/\kappa)$	$u^{(k)}$, $\bar{a}^{(k)}$, $\bar{\chi}^{(k)\dagger}$
`voldemand_all_open_only`	Net vega bought by non-market-makers, all traded options, trades that open a position; full-sample mode 0 on 60% of intervals	Cboe	none	$\text{asinh}(u/\kappa)$	$u^{(k)}$, $\bar{a}^{(k)}$, $\bar{\chi}^{(k)\dagger}$

continued on the next page

Table 6 (continued)

Variable	Description	Source	Diurnal adjustment	Transform $g(v)$	Predictors, $k \in \mathcal{K}$
FOMC (3; scheduled calendar)					
`fomc_day`	FOMC announcement-day dummy; full-sample mode 0 on 95% of intervals	FOMC calendar	none	v, not winsorized	$u^{(k)}$, $\bar{a}^{(k)}$, $\bar{\chi}^{(k)\dagger}$
`fomc_until_inv`	$1/(1+n^{+}_{t,i})$, $n^{+}_{t,i}$ intervals to the next announcement	FOMC calendar	none	$\log v$	$u^{(k)}$, $\bar{a}^{(k)}$
`fomc_since_inv`	$1/(1+n^{-}_{t,i})$, $n^{-}_{t,i}$ intervals since the last announcement	FOMC calendar	none	$\log v$	$u^{(k)}$, $\bar{a}^{(k)}$

Notes: Each series u enters as $u^{(k)}$, its adjusted, transformed and winsorized value averaged over the previous k intervals, and as $\bar{a}^{(k)}(u)$, the share of those intervals on which it was observed. †Closing-interval models only: $\bar{\chi}^{(k)}(u)$, the share of those intervals off the series' mode, for the 6 series whose full-sample mode accounts for at least 20% of intervals (mode and share in the description), the rule of the closing-interval forecasts reported, which predate the causal rule of Section 4. The pooled intraday models of Table 1 carry no $\bar{\chi}^{(k)}$. Under the causal rule, the specification, 5 series are mode-inflated at some interval: `numobs`, `stocktwits_attention`, `stocktwits_sentiment`, `stocktwits_sentcount` and `fomc_day`. $\mathcal{K} = \{1, 5, 25, 125, 625, 3125\}$ in the closing-interval model and in the intraday model; both models use every group here (Section 3). Diurnal adjustment: division by the trailing same-interval mean (mean) or standard deviation (s.d.); g: the variance-stabilizing map; the volatility-demand series take $\operatorname{asinh}(u_{t,i}/\kappa_{t,i})$ in place of both where the routing rule of Section 4 applies, on 99.4% or more of their intervals after the first print. A series is signed when the quantity it measures can be negative, not when its sample has a negative value. The columns give the chain of the pooled intraday forecasts of Table 1, in which the standard-deviation floor and the inverse hyperbolic sine routing use earlier intervals only. The closing-interval forecasts predate that chain: they take both over the full sample, call a series signed when any value in its full sample is negative, and log and winsorize the FOMC day dummy. In this table that changes two rows: `spread_vwstock`, signed by a single negative print (2018-11-12 11:00 ET) and so given $g(v) = v$ and the s.d. adjustment, and `fomc_day`, given $\log v$. Unnormalized and unscaled: without the factors $\sqrt{N_{t,i}}/\mathrm{RV}_{t,i}^{3/2}$ (skewness), $N_{t,i}/3$ (quarticity) and $\pi/2$ (bipower variation).

I Statistical tests

- **Diebold–Mariano test.** To compare the predictive accuracy of two competing forecasts A and B, we apply the Diebold and Mariano (1995) test to the daily QLIKE loss differentials, incorporating the small-sample adjustment of Harvey et al. (1997) for one-step-ahead forecasts:

$$d_t = L_t^A - L_t^B, \qquad \mathrm{DM} = \sqrt{\frac{|\mathcal{O}|-1}{|\mathcal{O}|}}\,\frac{\bar{d}}{\sqrt{\hat{\Omega}/|\mathcal{O}|}}.$$

  - Here A and B denote competing forecast series $f_{m,\mathcal{X},\mathcal{I}}$, indexed by the model class m, the predictor set $\mathcal{X} \in \{\mathrm{HAR}, \mathrm{Full}\}$, and the estimation interval set $\mathcal{I} \in \{\{c\}, \{1,\dots,M\}\}$.
  - The evaluation set $\mathcal{O}$, indexed chronologically by $t = 1,\dots,|\mathcal{O}|$, equals the full evaluation span $\mathcal{O} = \mathcal{T}_{\mathrm{oos}}$ ($|\mathcal{T}_{\mathrm{oos}}| = T_{\mathrm{oos}}$) for statistical accuracy comparisons and the option-data subsample $\mathcal{T}_{\mathrm{opt}} \subset \mathcal{T}_{\mathrm{oos}}$ for straddle return comparisons (Section 3).
  - We denote by $\bar{d}$ the sample mean of d_t over $\mathcal{O}$ and compute two-sided p-values under the standard normal distribution as $p = 2\{1 - \Phi(|\mathrm{DM}|)\}$.
  - Under this sign convention, $\mathrm{DM} < 0$ indicates that forecast A achieves a lower average loss than forecast B. Throughout the paper, "statistically significant" refers to significance at the 5% level ($p < 0.05$).
- **Newey–West long-run variance.** Following Newey and West (1987), we estimate the long-run variance $\hat{\Omega}$ using Bartlett kernel weights with lag truncation parameter q_{NW}:

$$\hat{\Omega} = \hat{\gamma}_0 + 2\sum_{j=1}^{q_{\mathrm{NW}}}\Big(1 - \frac{j}{q_{\mathrm{NW}}+1}\Big)\hat{\gamma}_j, \qquad \hat{\gamma}_j = \frac{1}{|\mathcal{O}|}\sum_{t=j+1}^{|\mathcal{O}|}(d_t - \bar{d})(d_{t-j} - \bar{d}).$$

E13. Decision: regenerate the remaining DM statistics.
Problem: `dm_test` now divides $\hat{\gamma}_j$ by $|\mathcal{O}|$, as displayed. The accuracy tables and macros were regenerated on 2026-10-18, but the DM statistics and significance stars in the appendix tuning and study tables are read from stored study outputs computed with the old divisor $|\mathcal{O}| - j$, so they predate the change and must be regenerated.

  - For the Diebold–Mariano test, we set the bandwidth according to the plug-in rule $q_{\mathrm{NW}} = \lfloor 4(|\mathcal{O}|/100)^{2/9} \rfloor$ (its numerical value on $\mathcal{T}_{\mathrm{oos}}$ is reported in Section 6.1).
- **Test of a mean daily return.** For any daily return series r_t—either the return of an individual trading strategy or the return differential Δ_t between two strategies—observed on the $T = 614$ post-warmup days $\mathcal{T}_{\mathrm{opt}}^{+}$ (Section 7.2), we test whether its mean is zero with the procedure that Lazarus et al. (2018) recommend for the Newey–West estimator:

$$t = \frac{\bar{r}}{\sqrt{\hat{\Omega}_r/T}}, \qquad \hat{\Omega}_r = \sum_{|j|<S}\Big(1 - \frac{|j|}{S}\Big)\hat{\gamma}_j, \qquad \hat{\gamma}_j = \frac{1}{T}\sum_{t=|j|+1}^{T}(r_t - \bar{r})(r_{t-|j|} - \bar{r}),$$

  with the Bartlett kernel, bandwidth $S = \lceil 1.3\,T^{1/2} \rceil = 33$, and fixed-$b$ critical values for $b = S/T$ (Kiefer and Vogelsang, 2005).
  - The autocovariances divide by T, with no small-sample factor; at lag truncation $S-1$ the estimator equals the Newey–West variance of the previous item.
  - This bandwidth is larger than that of conventional rules such as $0.75\,T^{1/3}$, which Lazarus et al. (2018) obtain from the formula of Andrews (1991) for a first-order autoregression with coefficient 0.25. Combined with fixed-b critical values, it reduces the size distortion of the test when the returns are serially correlated (Lazarus et al., 2018).

  - Under the null hypothesis and fixed b, the limiting distribution of t is a functional of Brownian motion that depends only on the kernel and on b (Kiefer and Vogelsang, 2005). We obtain its quantiles by simulation, computing the same statistic on 1,000,000 independent series of 1,000 i.i.d. standard normal draws with bandwidth $1{,}000\,b$. This gives the two-sided 5% critical value 2.11 and the p-values reported in the text, in place of the standard normal value 1.96.

- **Test of equal Sharpe ratios.** A test of $\bar{\Delta} = 0$ concerns the mean return differential, not the difference between two Sharpe ratios, which divides each strategy's mean return by its own standard deviation. For every statement about a Sharpe ratio difference $\Delta\mathrm{SR} = \mathrm{SR}(r^A) - \mathrm{SR}(r^B)$, with SR the ratio of Section 7.2, we therefore use a paired circular block bootstrap of $\Delta\mathrm{SR}$.
  - Each of 4,999 bootstrap samples joins blocks of 10 consecutive days, drawn with replacement from the T post-warmup days arranged on a circle; both strategies are resampled on the same days. Each drawn day keeps its own calendar year, and A_y stays at its value in the original sample, so each sample gives $\Delta\mathrm{SR}^*$ computed exactly as $\Delta\mathrm{SR}$.
  - The p-value is two-sided, the share of samples with $|\Delta\mathrm{SR}^* - \Delta\mathrm{SR}| \geq |\Delta\mathrm{SR}|$. The block length is fixed rather than calibrated, and the random seed is fixed.
  - The bootstrap is not studentized: SR averages ratios over calendar years, and no closed-form standard error of that average is used.
- **Test of equal Sharpe ratios of the pooled daily returns** (robustness). The studentized circular block bootstrap test of Ledoit and Wolf (2008) tests the equality of the Sharpe ratios of the daily returns pooled over the T days, the ratio of the mean to the standard deviation of all T returns. It is the test of the second alternative definition in Table 2 and is reported beside the bootstrap test above in Table 7.
  - With $u = (\mu_A, \mu_B, \mathrm{E}[(r_t^A)^2], \mathrm{E}[(r_t^B)^2])$, the difference of the Sharpe ratios of the daily returns, $\mu_A/\sigma_A - \mu_B/\sigma_B$, is a smooth function $f(u)$; its standard error follows from the delta method and the long-run covariance matrix Ψ of $(r_t^A, r_t^B, (r_t^A)^2, (r_t^B)^2)$.
  - The test statistic is $d = |f(\hat{u})|/\mathrm{se}(\hat{u})$. Each of 4,999 bootstrap samples joins blocks of 10 consecutive days, drawn with replacement from the T post-warmup days arranged on a circle, and gives $d^* = |f(\hat{u}^*) - f(\hat{u})|/\mathrm{se}(\hat{u}^*)$, with Ψ estimated in that sample by the block-based estimator of Ledoit and Wolf (2008). The p-value is $\big(\#\{d^* \geq d\} + 1\big)/(4{,}999 + 1)$.
  - In the original sample, Ψ is estimated with the Bartlett kernel at bandwidth equal to the block length, whose expectation is approximately that of the block-based estimator; Ledoit and Wolf (2008) use a prewhitened quadratic-spectral kernel at this step. The block length is fixed at 10 days rather than calibrated; the next item reports the number of rejections at block lengths 5 and 20. The random seed is fixed.
  - The test concerns the Sharpe ratios of daily returns over the T post-warmup days, without the year-by-year annualization of Section 7.2.
- **Robustness to the long-run variance estimator.** Table 7 lists every comparison of strategy returns whose statistical significance Section 7 or Appendix G.5 states. For each comparison it reports the t-statistic of the mean return differential with the Newey–West variance at lag truncation $q_{\mathrm{NW}} = \lfloor 4(T/100)^{2/9} \rfloor = 5$ (Newey and West, 1994) and at $q_{\mathrm{NW}} = \lfloor 1.5\,T^{1/3} \rfloor = 12$, the lag truncation of an earlier version of this paper, both with standard normal p-values; the fixed-b test above; the p-value of the bootstrap test of equal Sharpe ratios; and that of the Ledoit and Wolf (2008) test. Statements in the text about mean returns use the fixed-b test, and statements about Sharpe ratios use the bootstrap test.

  - Of the 34 tests of a mean, 0 are statistically significant at the 5% level under the fixed-b test, 0 at lag truncation 5 and 1 at lag truncation 12 with standard normal critical values.
  - Of the 33 tests of equal Sharpe ratios, 6 reject at the 5% level under the bootstrap test; under the Ledoit and Wolf (2008) test of the pooled ratios, 1 reject with block length 10, 1 with block length 5 and 1 with block length 20.
- **Multiple-comparison procedures.** For joint comparisons across predictor sets and model families, one may also consider the model confidence set of Hansen et al. (2011), the superior predictive ability test of Hansen (2005), and the StepM procedure of Romano and Wolf (2005).

  E14. Decision: drop the sentence, or report the set.
  Problem: the model confidence set, the superior predictive ability test, and StepM are named, and the paper reports none of their outputs.
  This bullet names three procedures and says they are for the comparison of predictor sets: the model confidence set, which returns the set of models not beaten at a stated confidence level; the superior predictive ability test, which returns a p-value against a benchmark; and StepM, which steps through the models and rejects some of them. No table in the paper states an MCS membership, an SPA p-value, or a StepM rejection. The next bullet says the main tables are point estimates and contain no bootstrap intervals, which is the kind of output those three procedures produce. The comparisons that are printed are Diebold–Mariano statistics on QLIKE and on squared error.
- **Reporting convention.** All main tables report point estimates accompanied by pairwise Diebold–Mariano statistics, t-statistics of mean returns with fixed-b critical values, or block-bootstrap p-values for Sharpe ratio differences, without bootstrap confidence intervals.

J Notes to the authors (not part of the paper)

- **Rerun pending (22 October 2026): the pooled numbers in this draft are preliminary.** Approved, and waiting for the window choice:
  - **Pooled intraday models** (Table 1, Sections 6 and 7): ridge, lasso, elastic net and OLS, LightGBM, XGBoost and the random forest, rerun under the final design: late-start zeroing (LATE_START_WARMUP), the neutral value for windows with no observation (UNOBSERVED_NEUTRAL), KKT-checked lasso and elastic-net walks (WALK_KKT_RESEED), the FOMC statement-interval dummy dropped (EXCLUDE_FOMC_RELEASE), and the rolling window chosen on the pre-2012 folds (window sweep running). The printed forecasts use a 1,000-day window without these changes (notes S4d, E28).
  - **Variable importance** (Section 6.5): the pooled linear results on the current design (results/rehaul_fi_pooled13_2026-10-22, preliminary) go in after the rerun; the scope of the tree analysis (SHAP by variable and state, interaction-free refits, SHAP interactions) is to be decided.
  - **Ablations:** leave one bucket out, then one variable out within the buckets that matter. Linear preview running; the final run is on the final design; trees to be decided.
  - **MinGRU network:** parked; rerun on the final design before it enters (tuning grid, seeds and refit cadence to be decided).
  - **Standardization** (STANDARDIZE_PENALTY, merged 2430543): the penalized pooled models are refit on columns standardized within each estimation window (Section 5.2). The window study now runs on a 2 January 2014 start, giving up 2012–13 of the accuracy sample, so that windows up to 3,500 days can be compared on the latest pre-sample fold; the first evaluation day of the final rerun follows that study.

Table 7: Tests of strategy returns under alternative long-run variance estimators

			NW, $q_{\mathrm{NW}} = 5$		NW, $q_{\mathrm{NW}} = 12$		NW, $S = 33$, fixed-b		Boot.	LW
	$\bar{\Delta}$ (%)	ΔSR	t	p	t	p	t	p	p	p
Panel A: Pooled intraday model minus short-straddle benchmark										
Ridge, HAR	6.61	1.19	1.38	0.168	1.35	0.178	1.35	0.198	0.125	0.175
Ridge, full	6.45	1.44	1.36	0.175	1.42	0.155	1.34	0.203	0.027	0.161
Lasso, HAR	6.49	1.17	1.35	0.176	1.32	0.187	1.33	0.208	0.131	0.185
Lasso, full	6.47	1.54	1.34	0.179	1.39	0.166	1.28	0.224	0.013	0.165
Elastic net, HAR	6.61	1.19	1.38	0.168	1.35	0.178	1.35	0.198	0.125	0.175
Elastic net, full	7.38	1.58	1.51	0.132	1.54	0.125	1.44	0.172	0.013	0.114
LightGBM, HAR	3.59	0.75	0.72	0.474	0.76	0.450	0.79	0.447	0.327	0.470
LightGBM, full	7.46	1.61	1.47	0.141	1.46	0.145	1.33	0.206	0.014	0.149
XGBoost, HAR	3.94	0.71	0.77	0.441	0.81	0.418	0.82	0.431	0.397	0.425
XGBoost, full	8.82	1.67	1.80	0.072	1.82	0.069	1.70	0.108	0.008	0.066
Random forest, HAR	5.92	1.03	1.16	0.247	1.12	0.263	1.12	0.284	0.180	0.254
Random forest, full	5.97	1.26	1.23	0.217	1.25	0.210	1.18	0.262	0.057	0.209
Panel B: Pooled intraday model, full predictor set minus HAR benchmark										
Ridge	−0.16	0.25	−0.05	0.961	−0.05	0.962	−0.05	0.962	0.657	0.963
Lasso	−0.02	0.37	−0.01	0.995	−0.01	0.995	−0.01	0.995	0.496	0.996
Elastic net	0.77	0.40	0.30	0.768	0.31	0.754	0.33	0.754	0.360	0.754
LightGBM	3.87	0.86	1.11	0.265	1.13	0.259	1.14	0.276	0.170	0.285
XGBoost	4.87	0.97	1.37	0.172	1.40	0.162	1.44	0.171	0.180	0.181
Random forest	0.05	0.23	0.02	0.987	0.02	0.987	0.02	0.986	0.739	0.985
Panel C: Pooled intraday minus closing-interval model										
Ridge, HAR	−2.58	−0.40	−0.84	0.400	−0.88	0.378	−0.82	0.433	0.370	0.366
Ridge, full	−1.82	−0.16	−0.55	0.581	−0.54	0.592	−0.50	0.634	0.698	0.599
Lasso, HAR	−2.71	−0.42	−0.88	0.377	−0.93	0.354	−0.86	0.410	0.350	0.342
Lasso, full	−3.62	−0.02	−0.98	0.325	−0.93	0.350	−0.85	0.415	0.964	0.350
Elastic net, HAR	−2.58	−0.40	−0.84	0.400	−0.88	0.378	−0.82	0.433	0.370	0.366
Elastic net, full	−3.86	−0.21	−1.20	0.231	−1.26	0.209	−1.15	0.271	0.536	0.220
LightGBM, HAR	2.19	0.32	0.53	0.596	0.51	0.613	0.51	0.627	0.637	0.643
LightGBM, full	0.52	0.31	0.10	0.922	0.09	0.925	0.09	0.933	0.699	0.922
XGBoost, HAR	0.58	−0.16	0.21	0.835	0.21	0.833	0.24	0.819	0.711	0.845
XGBoost, full	2.97	0.47	0.61	0.543	0.60	0.551	0.52	0.620	0.419	0.554
Random forest, HAR	4.94	0.66	1.09	0.278	1.08	0.282	1.10	0.293	0.319	0.287
Random forest, full	2.41	0.30	0.68	0.499	0.67	0.503	0.60	0.567	0.550	0.506
Panel D: Closing-interval model, full predictor set minus HAR benchmark										
Elastic net	2.05	0.21	0.78	0.436	0.78	0.437	0.74	0.476	0.607	0.450
Lasso	0.89	−0.03	0.35	0.727	0.34	0.733	0.35	0.739	0.942	0.737
Panel E: Closing-interval model minus short-straddle benchmark										
Elastic net, full	11.24	1.80	1.95	0.051	2.00	0.046	1.92	0.072	0.018	0.035
Panel F: Short-straddle benchmark, mean return										
Short straddle	2.43	–	0.58	0.559	0.62	0.536	0.67	0.522	–	–

Notes: Each row is a comparison whose statistical significance the text states, on the 614 post-warmup days of $\mathcal{T}_{\mathrm{opt}}$. $\bar{\Delta}$ is the mean daily return differential at mid-quotes in percent of the straddle premium (the mean return itself in the last panel), and ΔSR the difference of the annualized Sharpe ratios of Section 7.2. The t-statistics test $\bar{\Delta} = 0$ with the Newey–West long-run variance: at lag truncation $q_{\mathrm{NW}} = \lfloor 4(T/100)^{2/9} \rfloor = 5$ (Newey and West, 1994) and at $q_{\mathrm{NW}} = \lfloor 1.5\,T^{1/3} \rfloor = 12$, both with standard normal p-values, and at bandwidth $S = \lceil 1.3\,T^{1/2} \rceil = 33$ with fixed-b p-values (Lazarus et al., 2018; Kiefer and Vogelsang, 2005), $T = 614$. Boot. is the p-value of the paired block-bootstrap test of equal Sharpe ratios SR (Section 7.2; block length 10, 4,999 resamples), and LW that of the Ledoit and Wolf (2008) test of equal Sharpe ratios of the daily returns pooled over the T days (block length 10, 4,999 resamples). The text uses the fixed-b test for statements about mean returns and the bootstrap test for statements about Sharpe ratios. All p-values are two-sided.

  - **Closing-interval models:** not rerun (James, 22 October 2026); their full-sample rules are disclosed (notes S4a, S4c, S4d). The data fixes they predate (invalid-print nulling, FOMC calendar, VIX splice, fill cap) still need a disclosure in Section 3 (E28).
  - **Predictor count:** 586 still counts the calendar and interaction columns cut by note E9 and the FOMC statement dummy; the printed design has 576 columns and the rerun's will have 564. Take the count from the design at the rerun.
- **Accuracy against the benchmark:** 0 of the 12 pooled intraday forecasts have lower QLIKE than the daily HAR benchmark, 0 significantly (Diebold–Mariano, 5% level); closing-interval tests are left out of this count (Appendix G).
- **Data issue to fix before final numbers:**
  - the activity flags are computed after the forward fill, so the FOMC flags read 1 on all 122 trade days after November 2023;

    S4c. Decision: no rerun of the closing-interval models (James, 22 October 2026; note S4a): the flags read after the fill stay in the printed closing-interval forecasts, disclosed in Section 4 (activity indicators) and Table 6.
    Problem: this item is not marked as true of the closing-interval forecasts only, and it is no longer a fix to make before final numbers. The pooled intraday forecasts of Table 1 were built after commit b29774a (CAUSAL_MODE_INFLATION and ACTIVE_FROM_OBSERVED in src/backtest/executor.py: flags from observed prints) and carry no activity averages (note E6). The code comment counts 123 FOMC days, not 122.
  - VIX, volume-demand and FOMC feeds end early and are carried forward without limit into 2024.
- **Missing outputs:** paired-comparison tables (axis 1, axis 2), DM against the OLS in the master table, and the model confidence set over models.
- **Pending:** the pooled random forest (full table).
- **E15, sample and window still open** (2026-10-16). Hyperparameters: one configuration per model, chosen once on the pre-sample and frozen on the evaluation span. The daily HAR OLS stays outside the scheme. Not applied to Table 1.
  - Open choice: all 48 intervals, overnight included, or the 13 regular-hours bars (end stamps 10:00–16:00). Window: 24,000 intervals or 96,000 intervals (2,000 trading days on the 48-bar grid).
  - Calendar (note E9): columns after the overnight indicator are out of the specification, including the products of the target HAR components with the opening- and closing-interval indicators. The overnight indicator stays only if the overnight intervals stay in the training sample. On the 13 regular-hours bars it equals 1 only on the 16:00 bar. Not applied to the printed column counts or Table 1.
  - Shorter window, 16:00 rows only (closing-interval window study, 2026-10-03; trees refit every 10 sessions and the linear penalties are re-chosen, so this is not the frozen rule). LightGBM QLIKE 0.1110 at 24,000/48 sessions and 0.1006 at 2,000 sessions (DM 2.71). Ridge 0.1194 and 0.1027 (DM 2.94). Lasso 0.0998 and 0.1006 (DM −0.36).
  - A 48-bar pool against a one-bar fit, HAR plus calendar, 2012–2024. The pool preferred the shorter window; the one-bar fits preferred 2,000 days. At the 16:00 bar the one-bar QLIKE gain was 6–8% on the shorter window and 12–13% at 2,000 days. The full predictor set on one bar was worse on the shorter window (+4.1%, 8 bars worse) and flat at 2,000 days (+0.5%).
  - The 13 regular-hours bars as one model against the 48-bar pool, 2,000 days, mean over buckets. QLIKE lower by 1.75% / 1.86% / 1.85% (ridge / elastic net / lasso); at the

close, 5.5–6.0%. The 15:30 sign trade, crossed, matched the pool (0.70 / 0.62 / 0.68 against the pool's 0.72 / 0.68 / 0.73). It moved only for a fit on the 16:00 bar alone (QLIKE 8.6–9.5% below the pool; crossed Sharpe 1.00 / 0.92 / 0.85, lasso / elastic net / ridge).

- **The full outline,** with source and macro pointers, is `writeup/PAPER_REHAUL_OUTLINE_2026-10-15.md`.